

# CS 106: Artificial Intelligence

## Constraint Satisfaction Problems

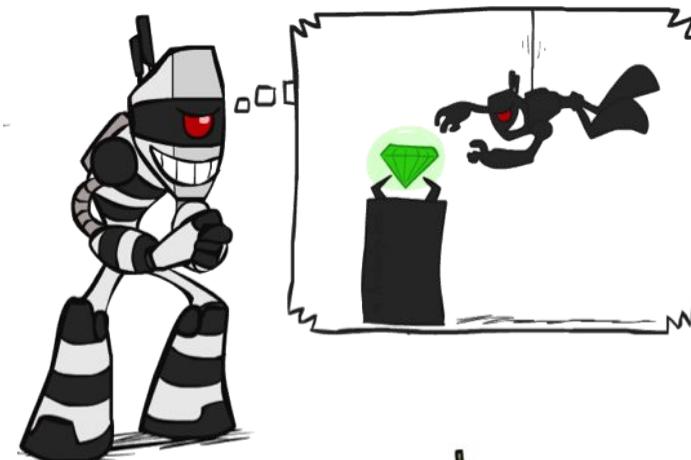


**Instructor:** Ngoc-Hoang LUONG, PhD  
**University of Information Technology (UIT), VNU-HCM**

[These slides were adapted from the slides created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley. <http://ai.berkeley.edu/>.]

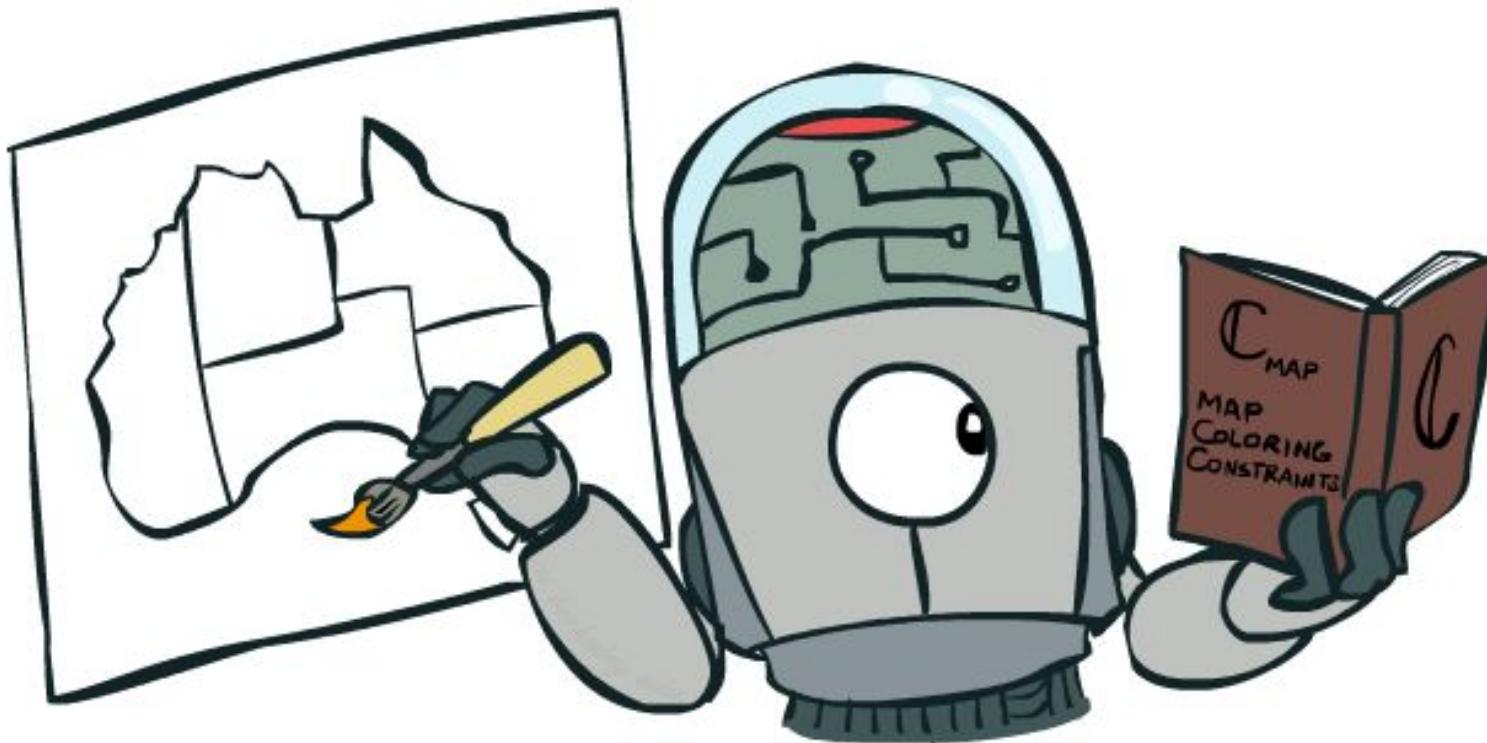
# What is Search For?

- Assumptions about the world: a single agent, deterministic actions, fully observed state, discrete state space
- Planning: sequences of actions
  - The path to the goal is the important thing
  - Paths have various costs, depths
  - Heuristics give problem-specific guidance
- Identification: assignments to variables
  - The goal itself is important, not the path
  - All paths at the same depth (for some formulations)
  - CSPs are specialized for identification problems



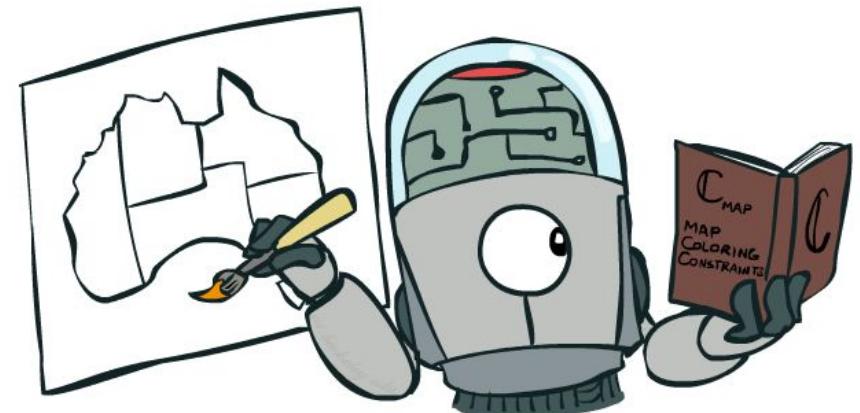
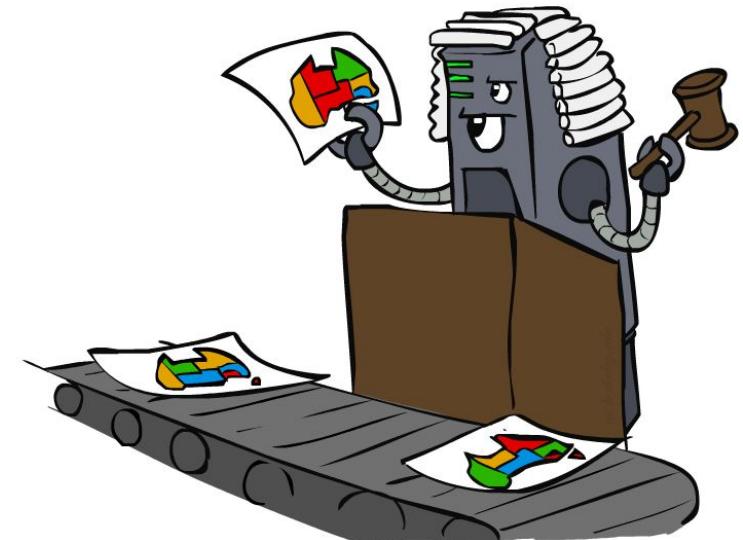
# Constraint Satisfaction Problems

---



# Constraint Satisfaction Problems

- Constraint satisfaction problems (CSPs):
  - A special subset of search problems
  - State is defined by **variables  $X_i$** , with values from a **domain  $D$**  (sometimes  $D$  depends on  $i$ )
  - Goal test is a **set of constraints** specifying allowable combinations of values for subsets of variables
- Allows useful general-purpose algorithms with more power than standard search algorithms

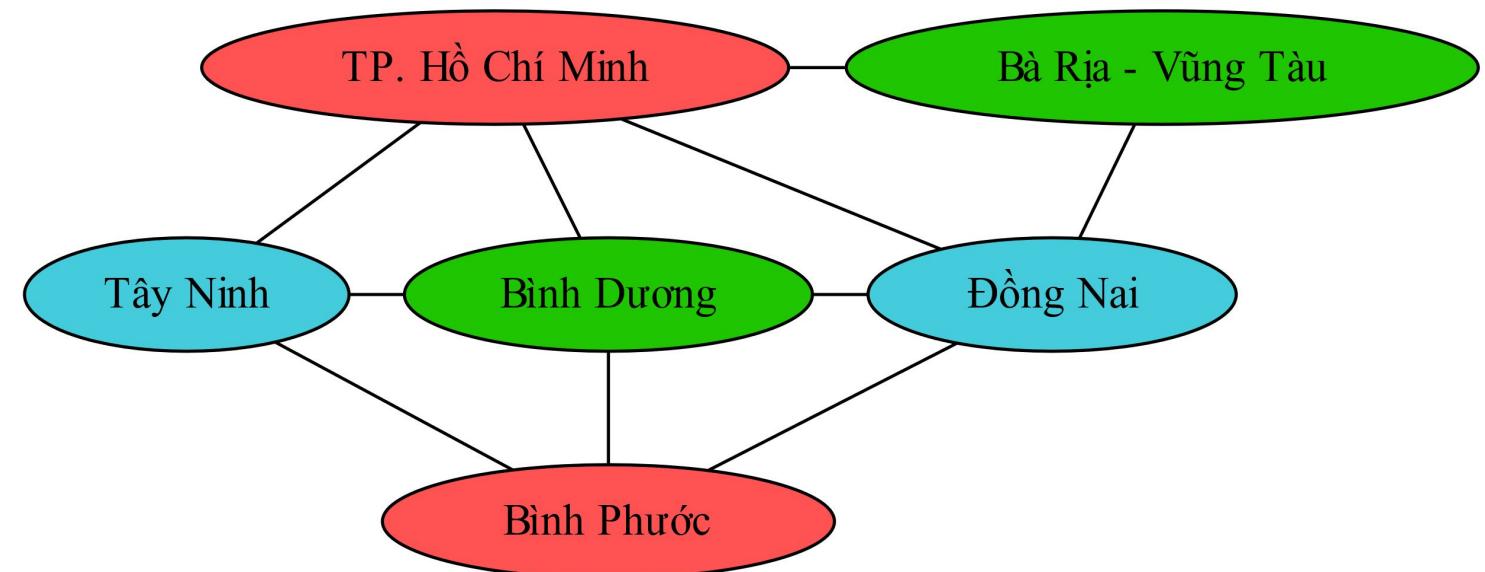


# Constraint Satisfaction Problems

---

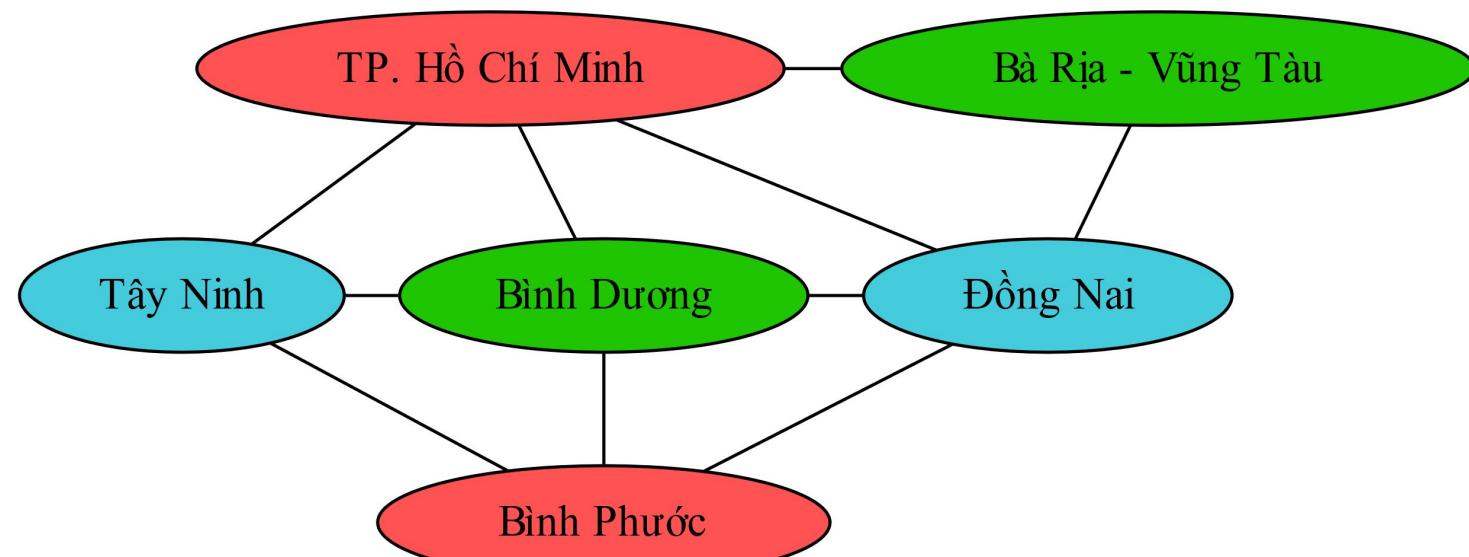
- A constraint satisfaction problem (CSP) is a tuple  $(X, D, C)$  where:
  - $X = \{x_1, x_2, \dots, x_n\}$  is the set of **variables**.
  - $D = \{d_1, d_2, \dots, d_n\}$  is the set of **domains**.
  - $C = \{c_1, c_2, \dots, c_m\}$  is a set of **constraints**.
- For example,  $x, y, z \in \{0,1\}, x + y = z$  is a CSP where:
  - Variables are:  $x, y, z$
  - Domains are:  $d_x = d_y = d_z = \{0,1\}$
  - There is a single constraint:  $x + y = z$

# Constraint Graphs

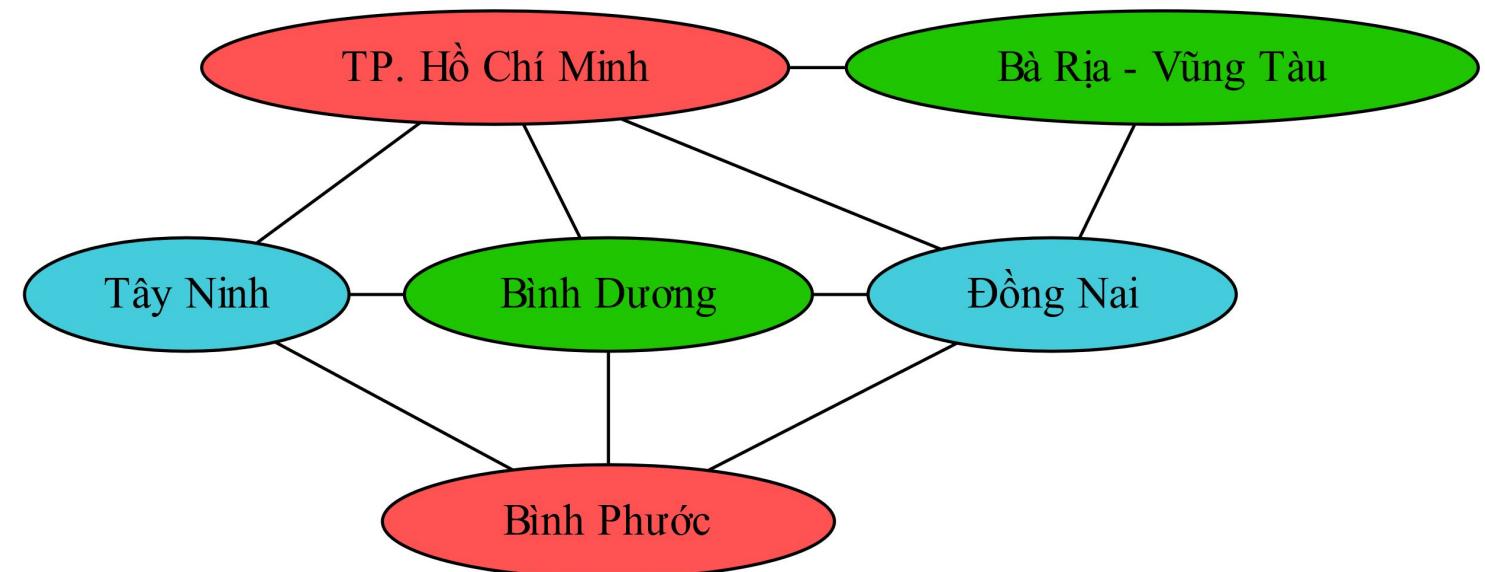


# Constraint Graphs

- Binary CSP: each constraint relates (at most) two variables
- Binary constraint graph: nodes are variables, arcs show constraints
- General-purpose CSP algorithms use the graph structure to speed up search.



# CSP Examples



# Example: Map Coloring

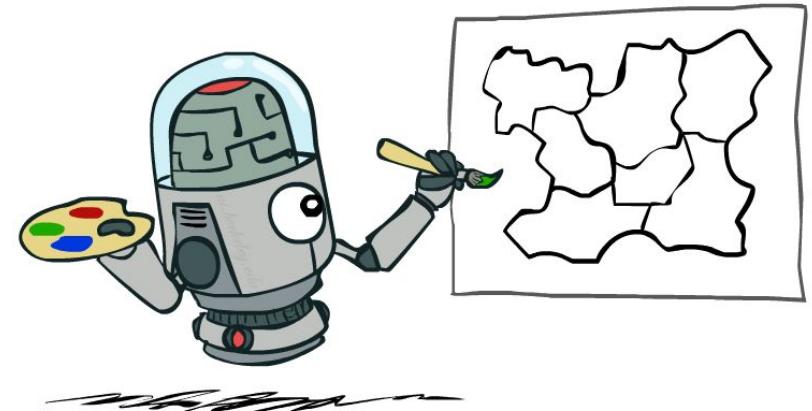
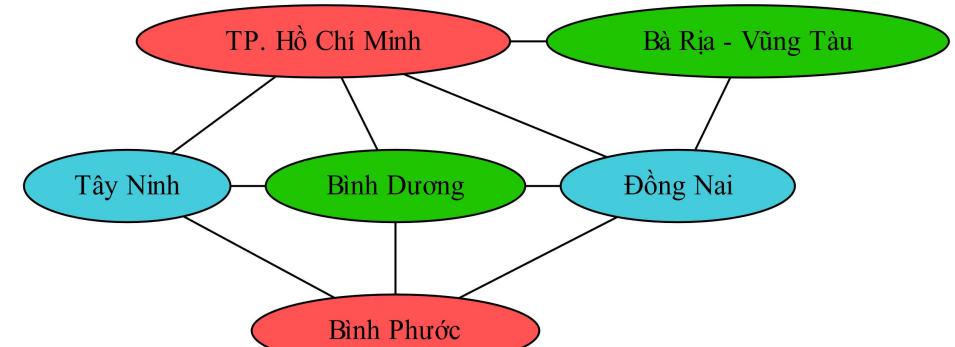
- Variables:  $\{HCM, TN, BD, BP, \bar{DN}, BR\}$
- Domains: {red, green, blue}
- Constraints: adjacent regions must have different colors

Implicit:  $HCM \neq TN$

Explicit:  $(HCM, TN) \in \{(red, green), (red, blue), \dots\}$

- Solutions are assignments satisfying all constraints, e.g.:

$$\{HCM = red, TN = blue, BD = green, \\ BP = red, \bar{DN} = blue, BR = green\}$$



# Example: N-Queens

- Variables:  $Q_k$
- Domains:  $\{1, 2, 3, \dots, N\}$
- Constraints:



Implicit:  $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

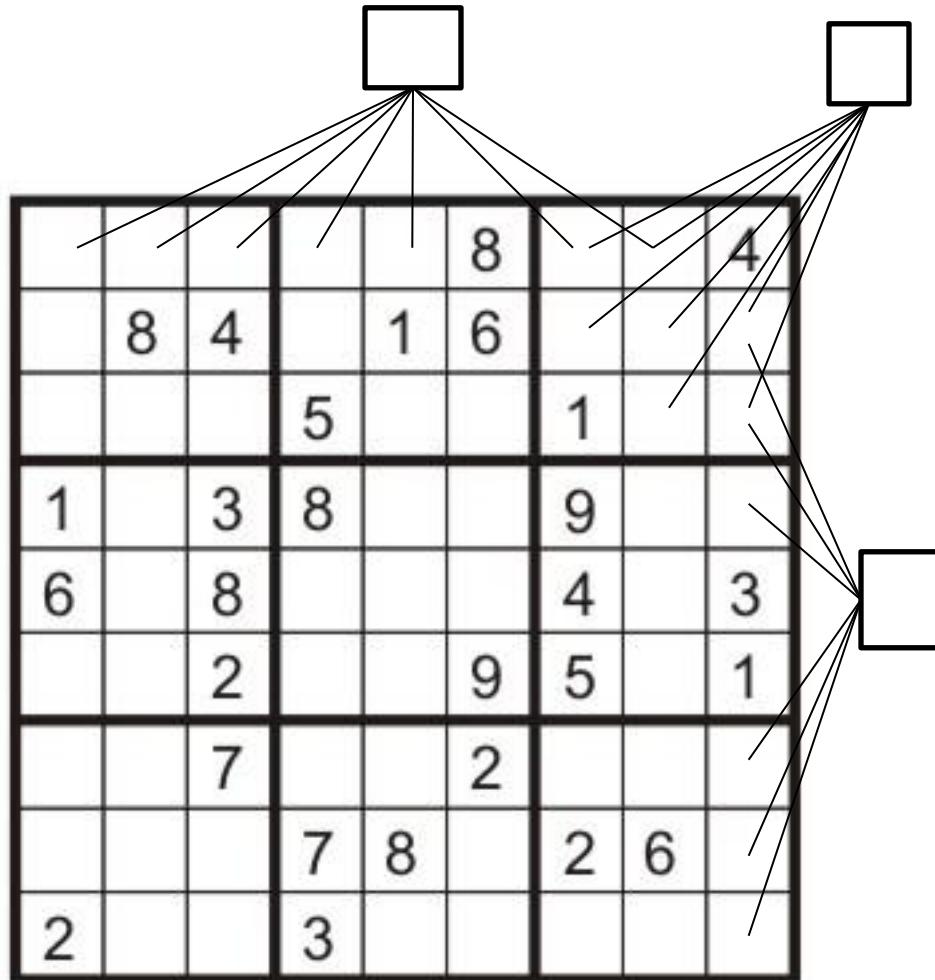
Explicit:  $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$

...

$Q_1$	$Q_2$	$Q_3$	$Q_4$
		Q	
Q			
			Q
	Q		

Q			
	Q		

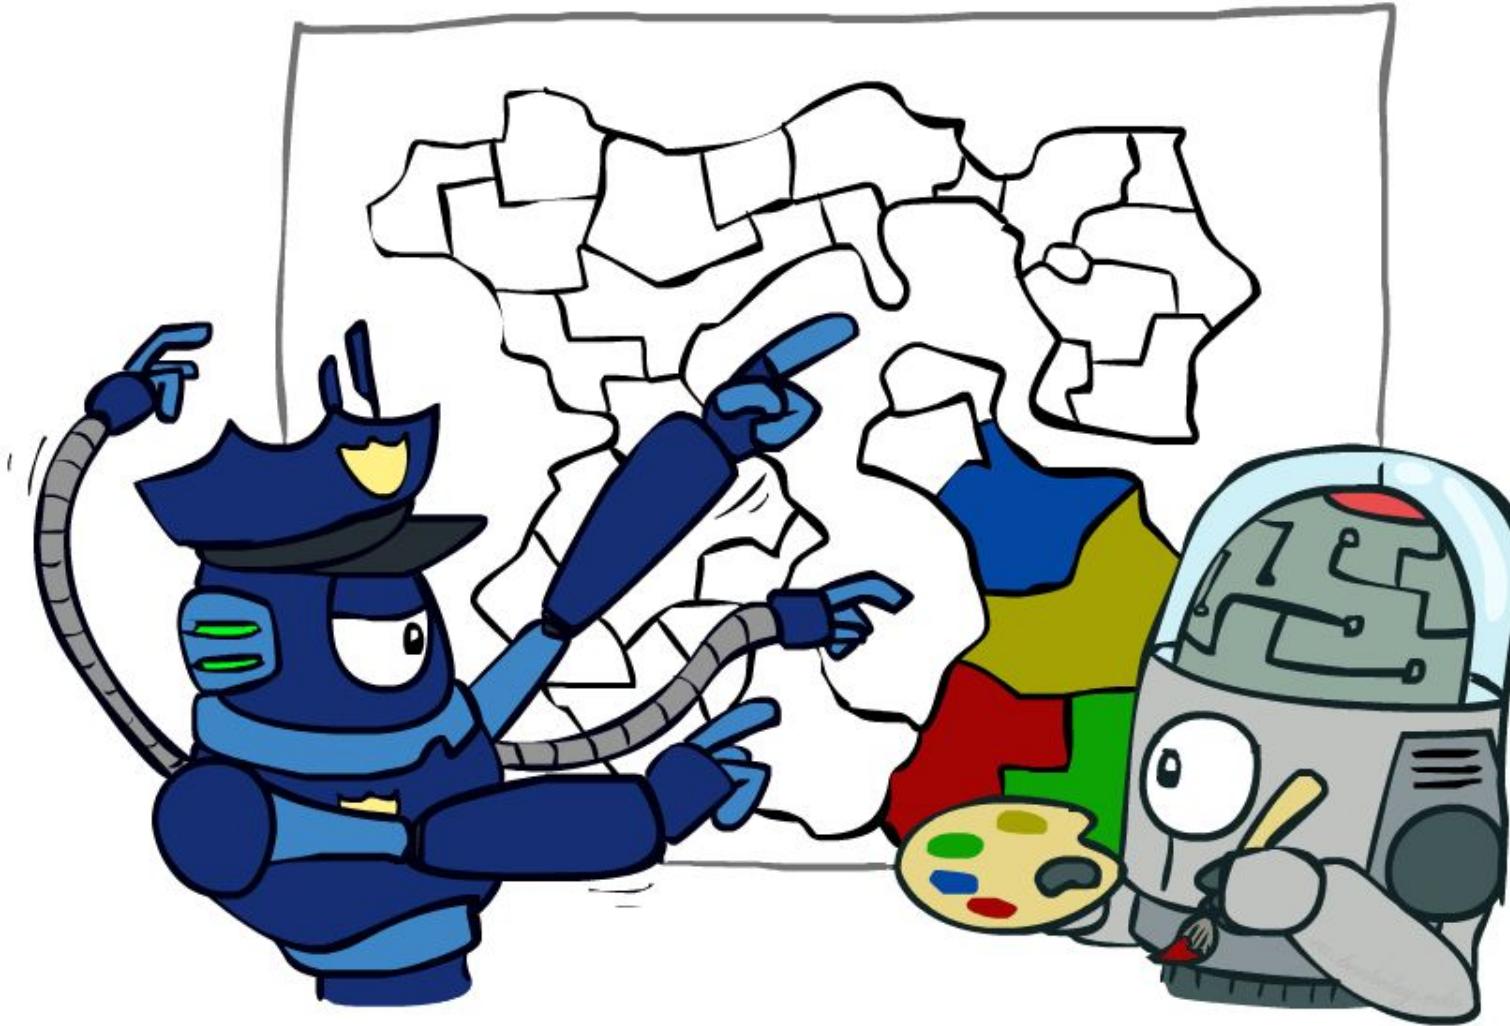
# Example: Sudoku



- Variables:
  - Each (open) square
- Domains:
  - $\{1, 2, \dots, 9\}$
- Constraints:
  - 9-way alldiff for each column
  - 9-way alldiff for each row
  - 9-way alldiff for each region

# Varieties of CSPs and Constraints

---



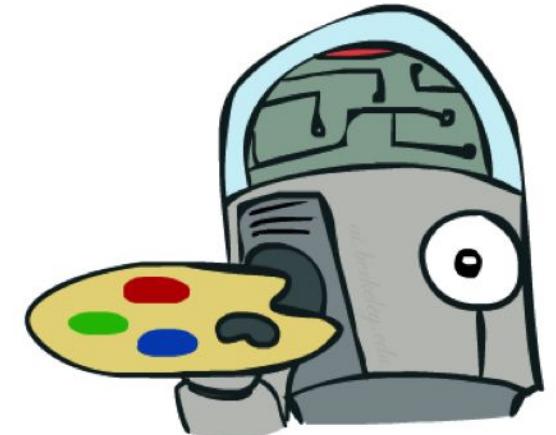
# Varieties of CSPs

- Discrete Variables

- Finite domains
  - Size  $d$  means  $O(d^n)$  complete assignments
  - E.g., Boolean CSPs, including Boolean satisfiability (NP-complete)
- Infinite domains (integers, strings, etc.)
  - E.g., job scheduling, variables are start/end times for each job

- Continuous variables

- E.g., start/end times for Hubble Telescope observations



# Varieties of Constraints

- Varieties of Constraints

- Unary constraints involve a single variable (equivalent to reducing domains), e.g.:

*TP.HCM ≠ blue*

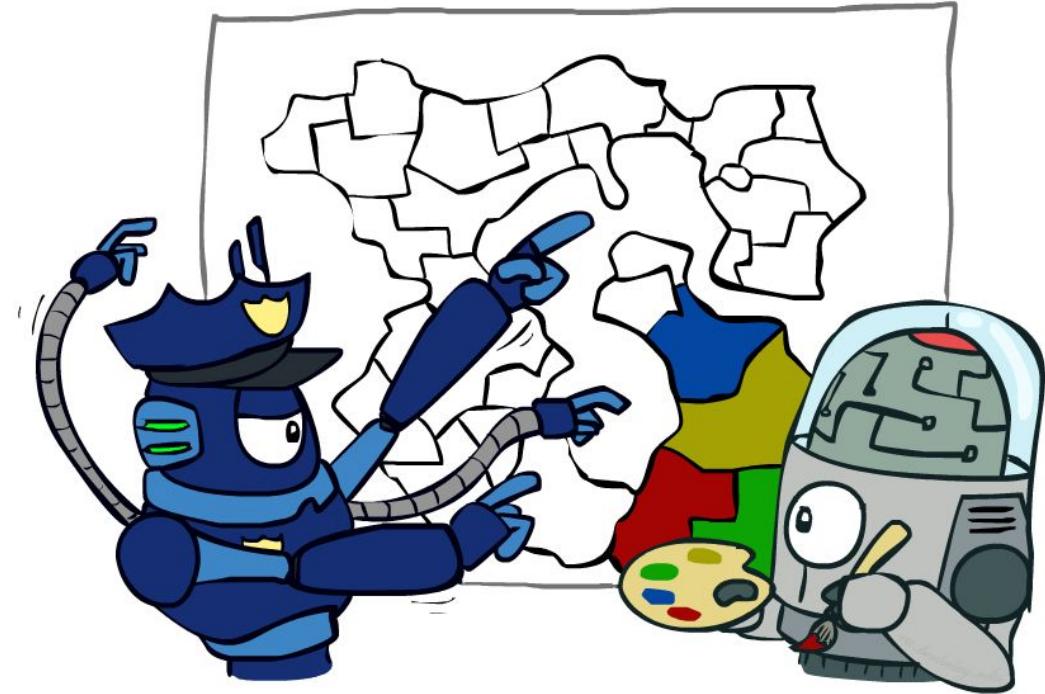
- Binary constraints involve pairs of variables, e.g.:

*TP.HCM ≠ Tay Ninh*

- Higher-order constraints involve 3 or more variables:  
e.g., Sudoku constraints

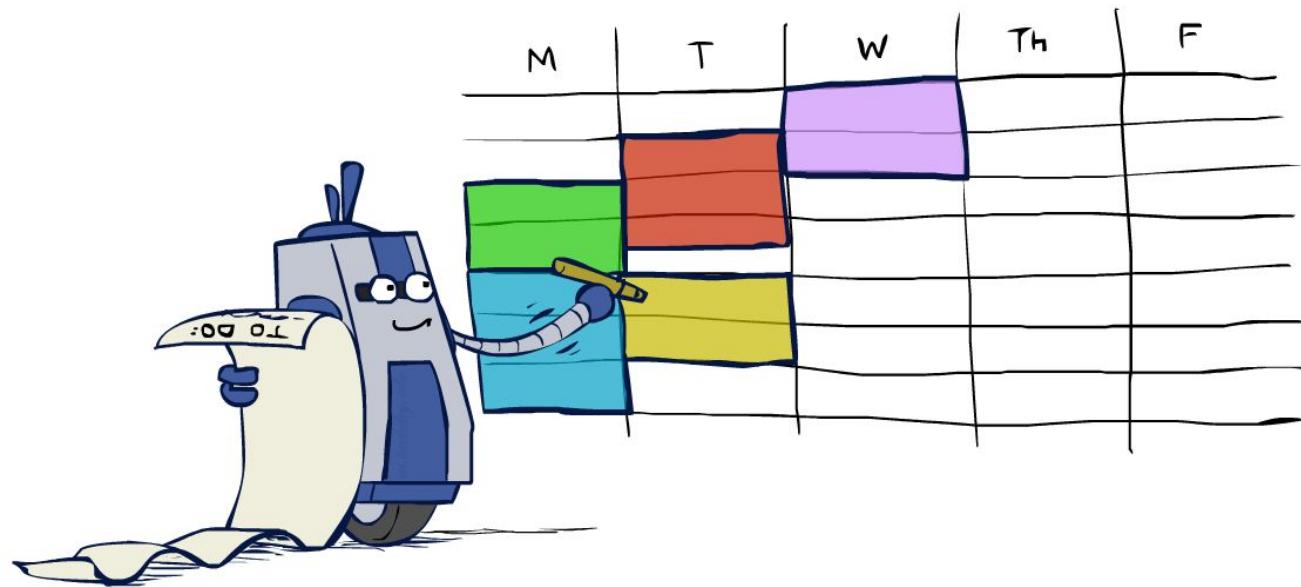
- Preferences (soft constraints):

- E.g., red is better than green
  - Often representable by a cost for each variable assignment
  - Gives constrained optimization problems



# Real-World CSPs

- Assignment problems: e.g., who teaches what class
- Timetabling problems: e.g., which class is offered when and where?
- Hardware configuration
- Transportation scheduling
- Factory scheduling
- Circuit layout
- Fault diagnosis
- ... lots more!



- Many real-world problems involve real-valued variables...

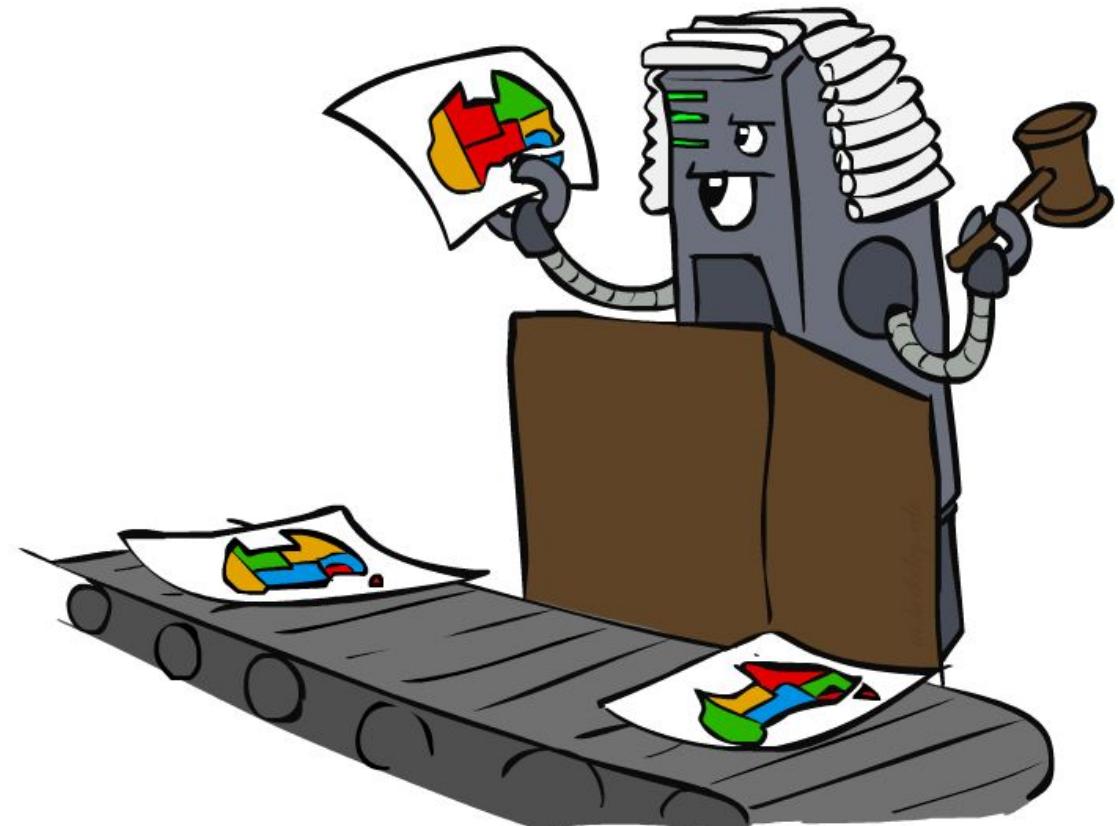
# Solving CSPs

---



# Standard Search Formulation

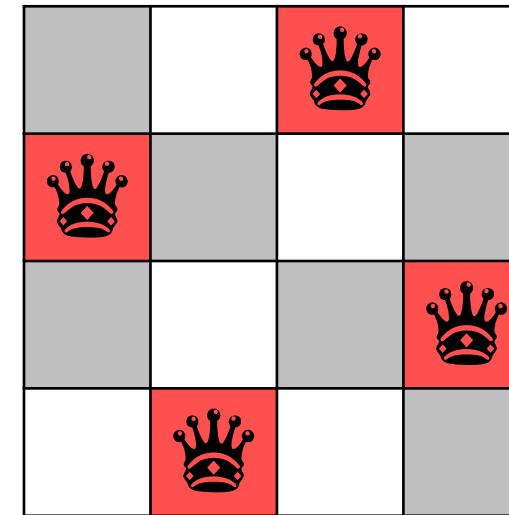
- Standard search formulation of CSPs
- States defined by the values assigned so far (**partial assignments**)
  - Initial state: the empty assignment, {}
  - Successor function: assign a value to an unassigned variable
  - Goal test: the current assignment is complete and satisfies all constraints
- We'll start with the straightforward, naïve approach, then improve it



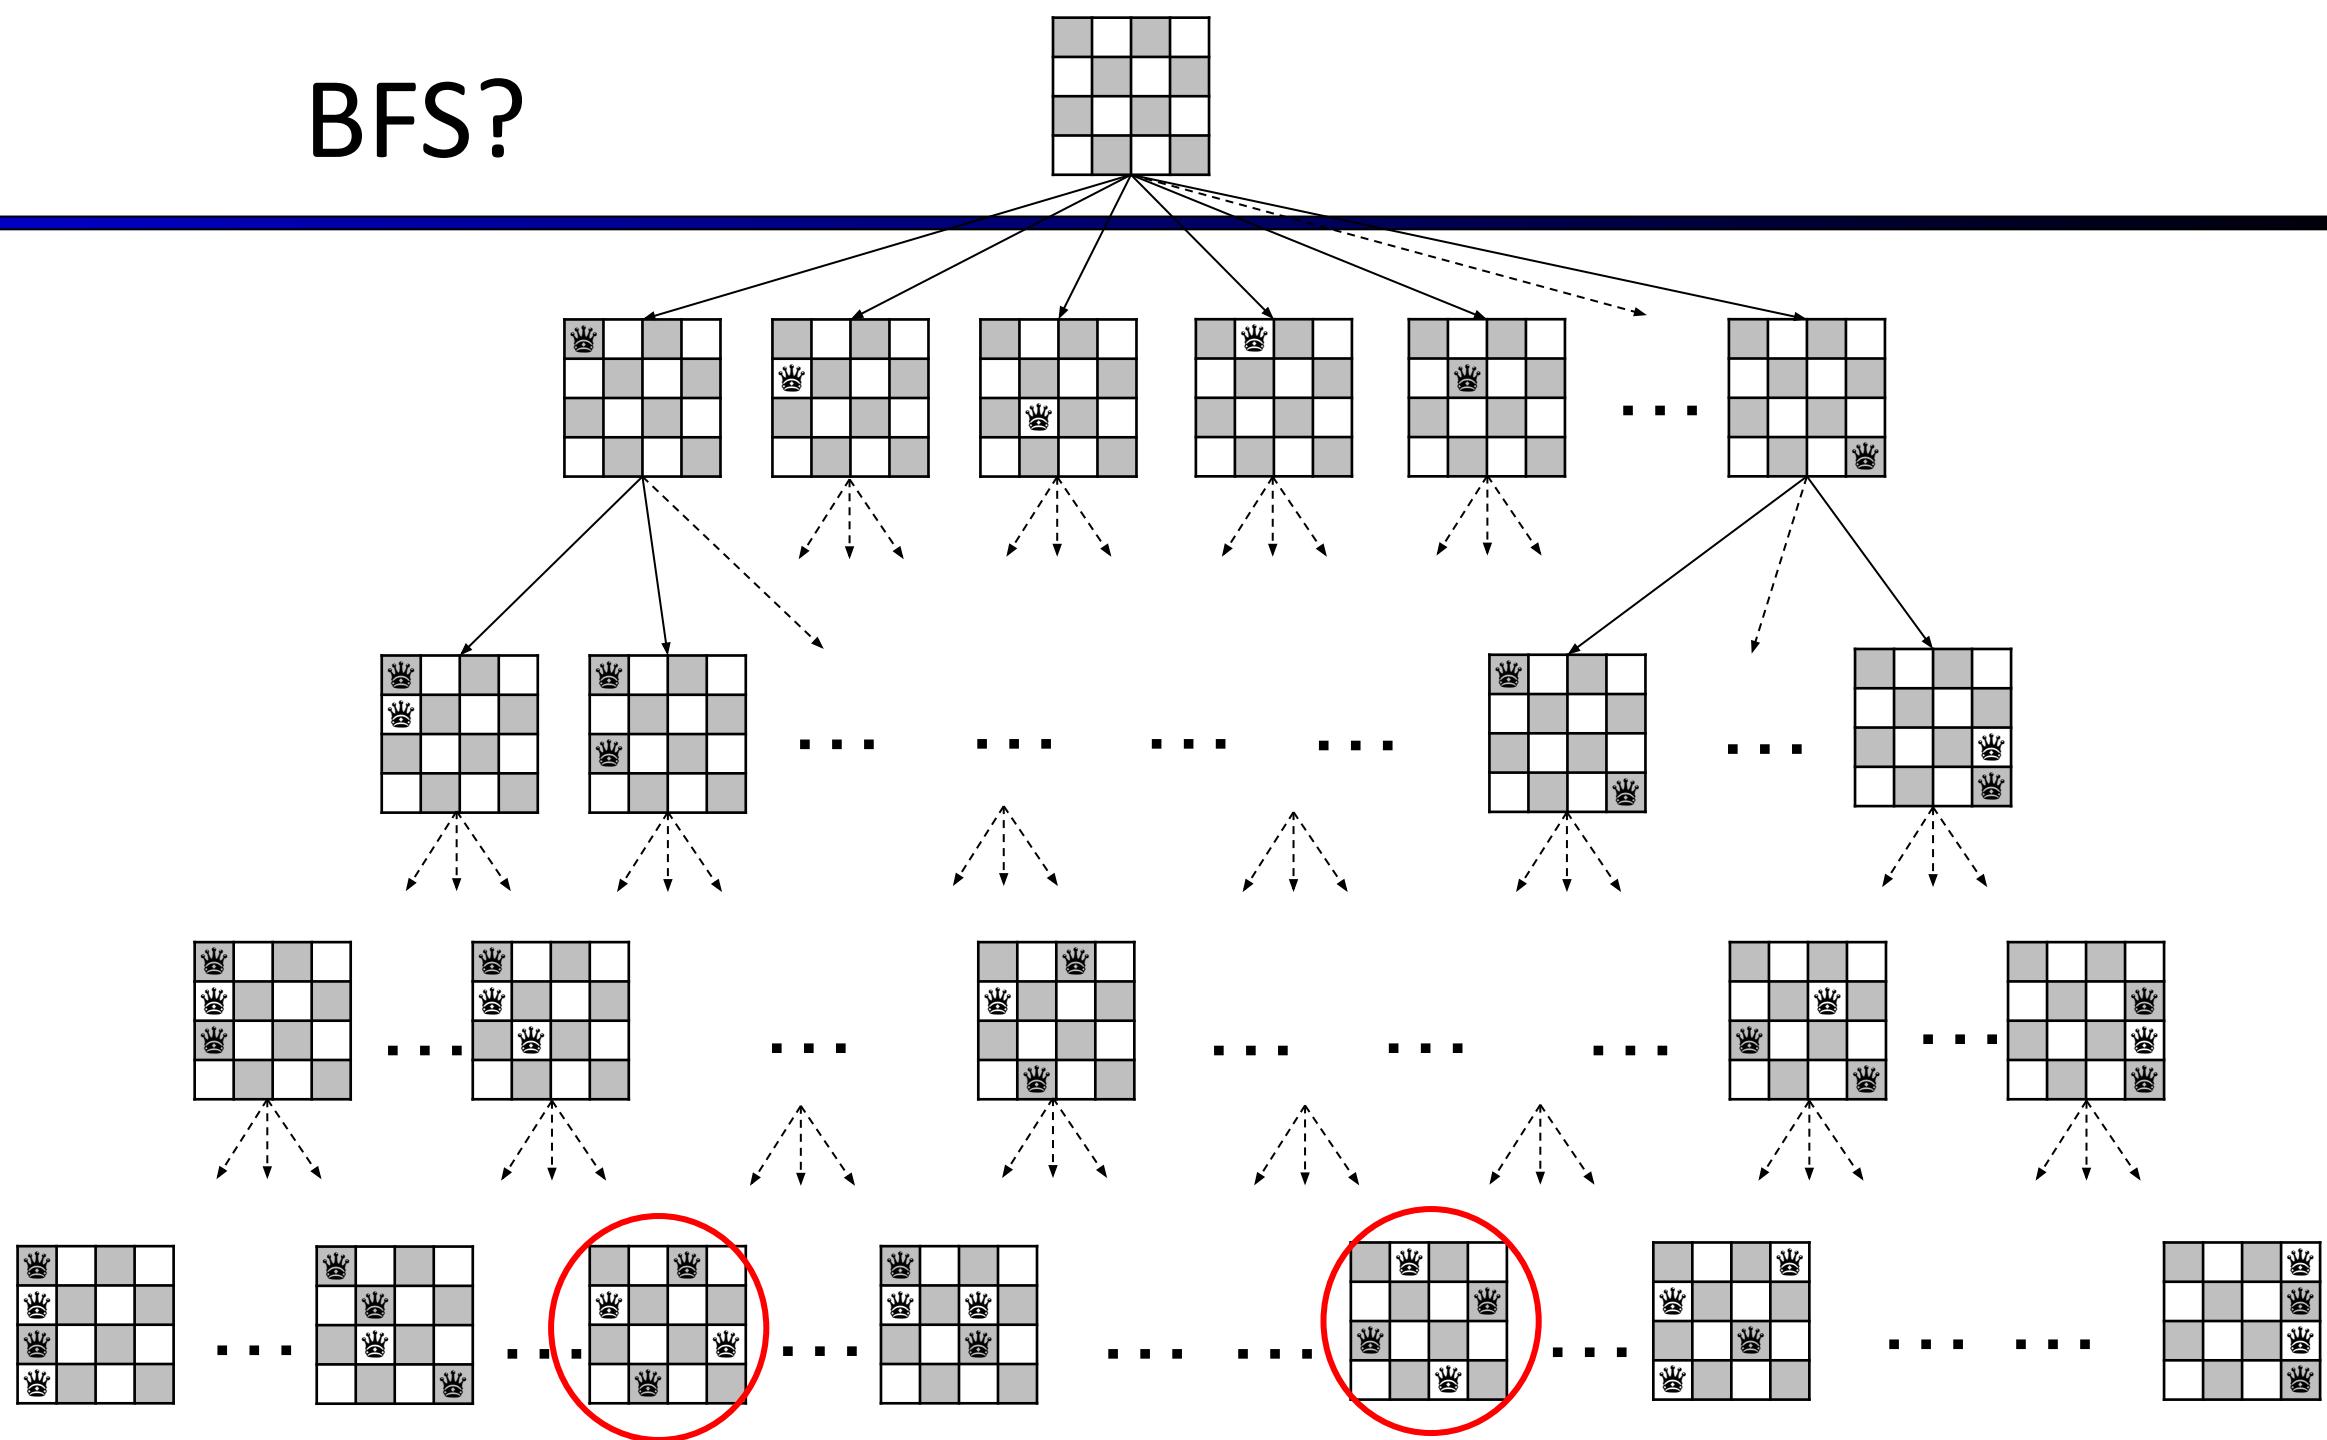
# Search Methods

---

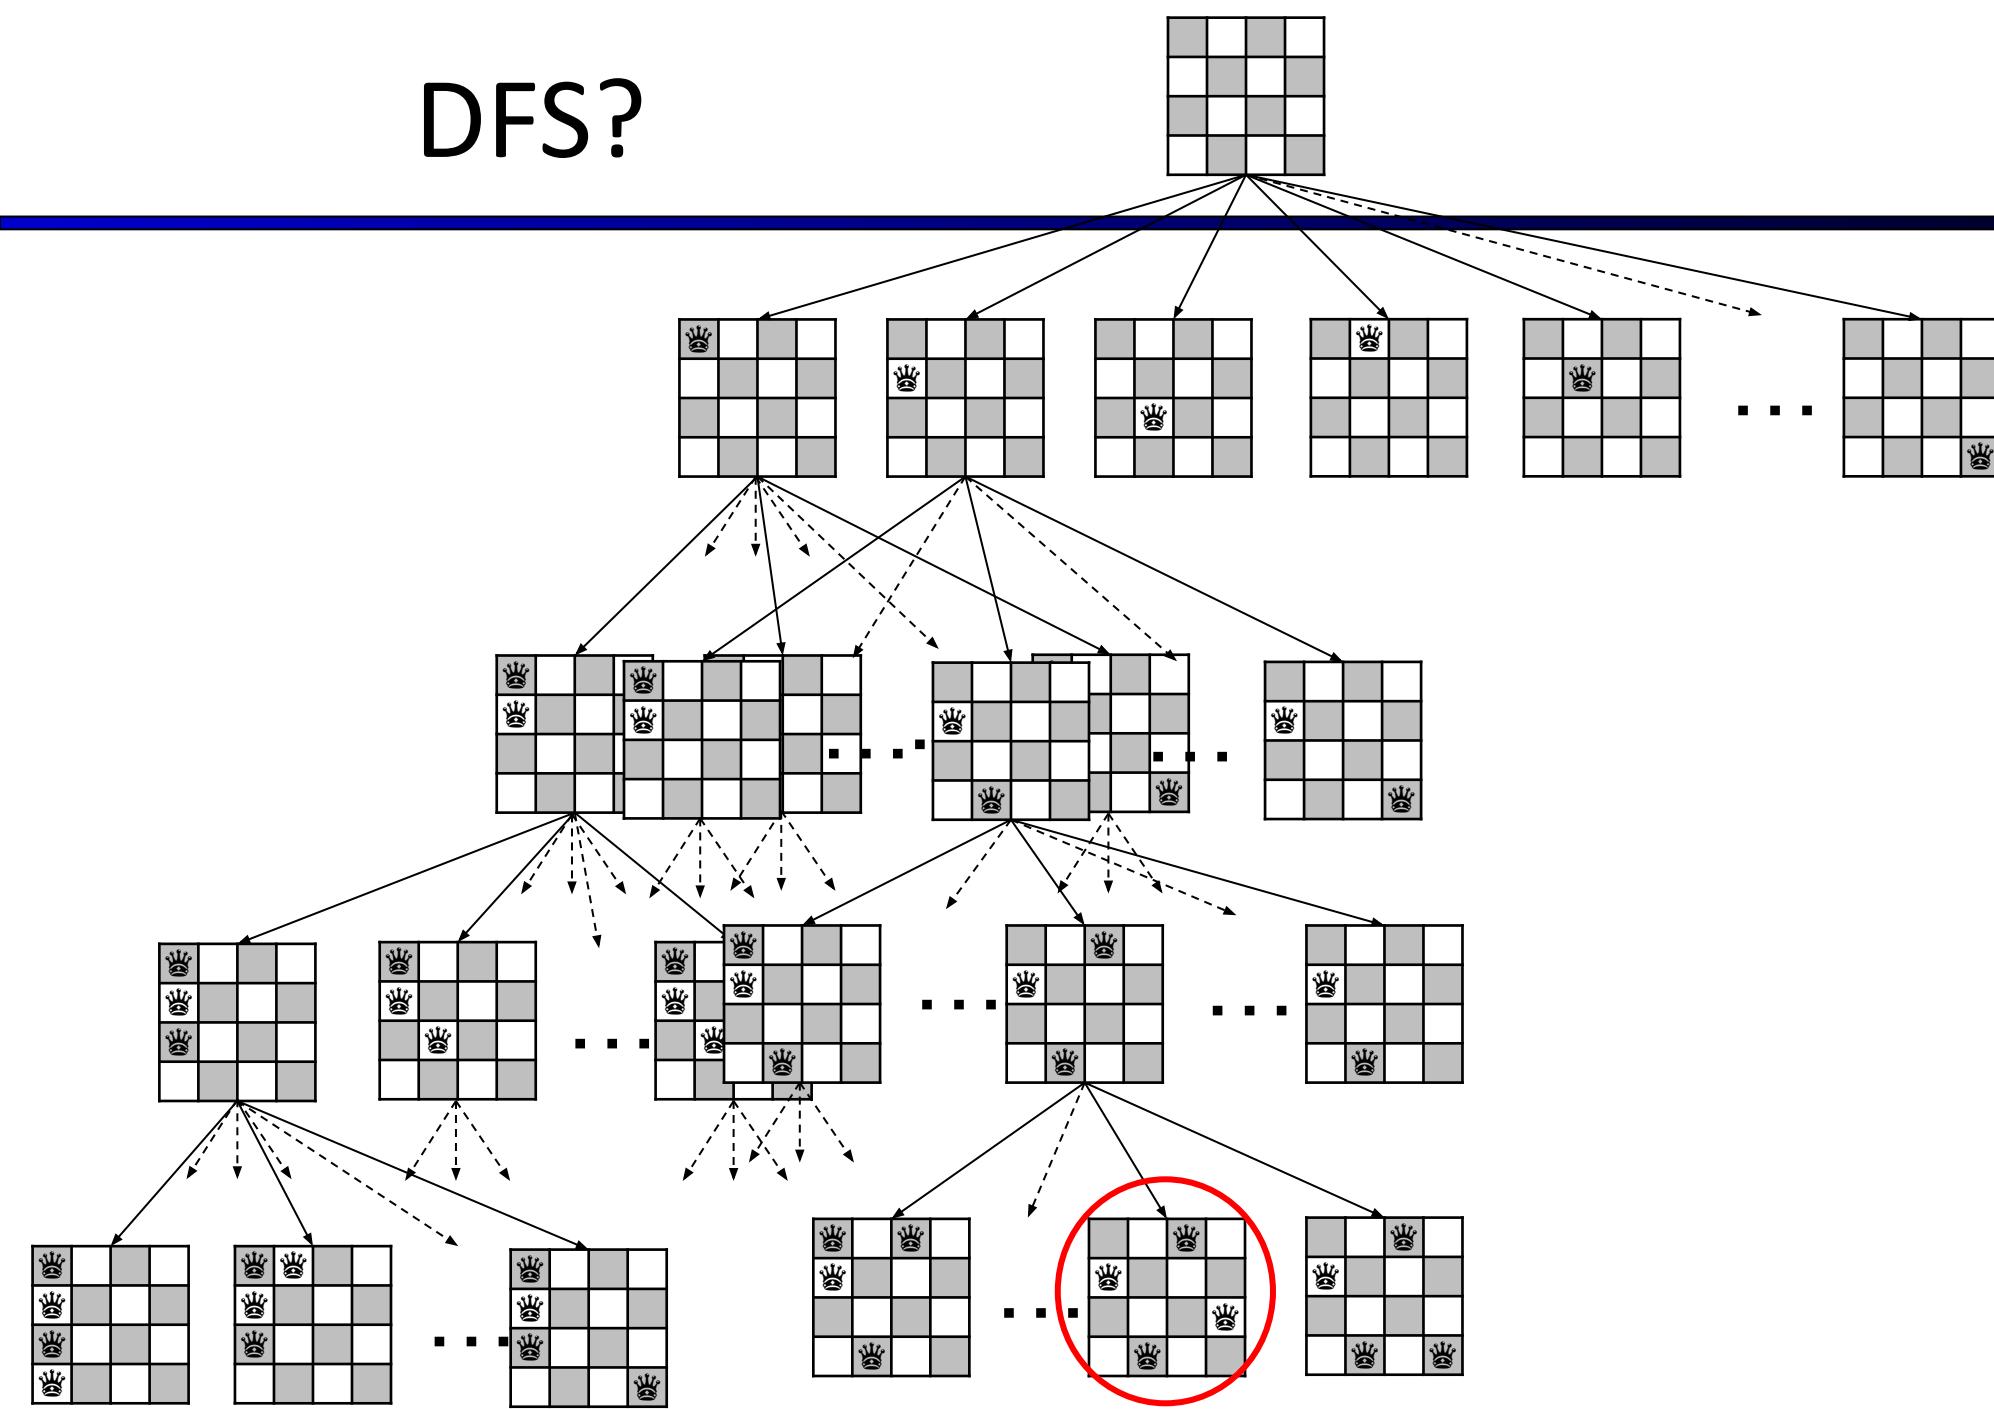
- What would BFS do?
- What would DFS do?
- What problems does naïve search have?



# BFS?

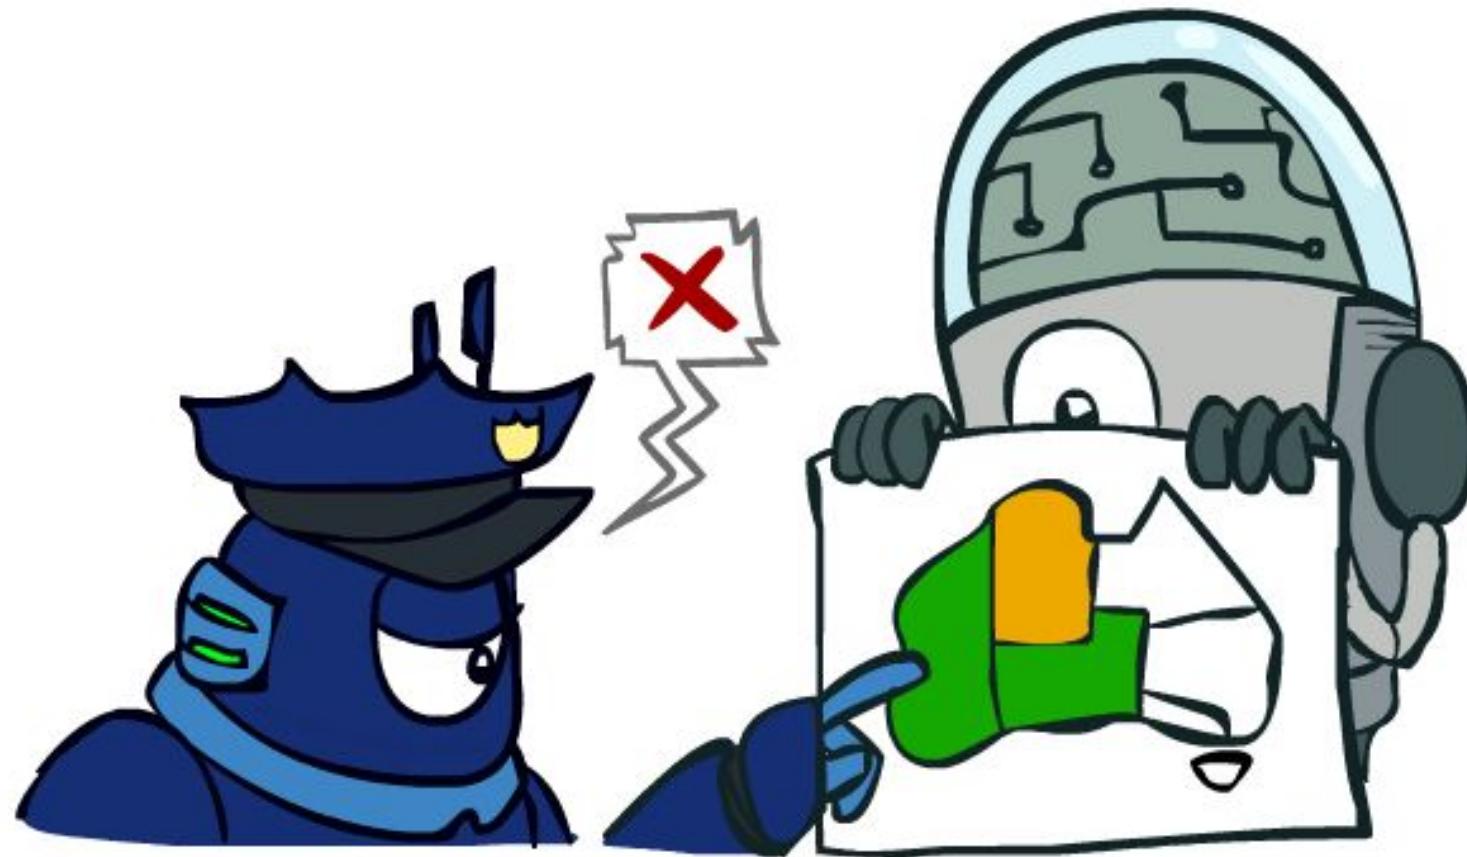


# DFS?



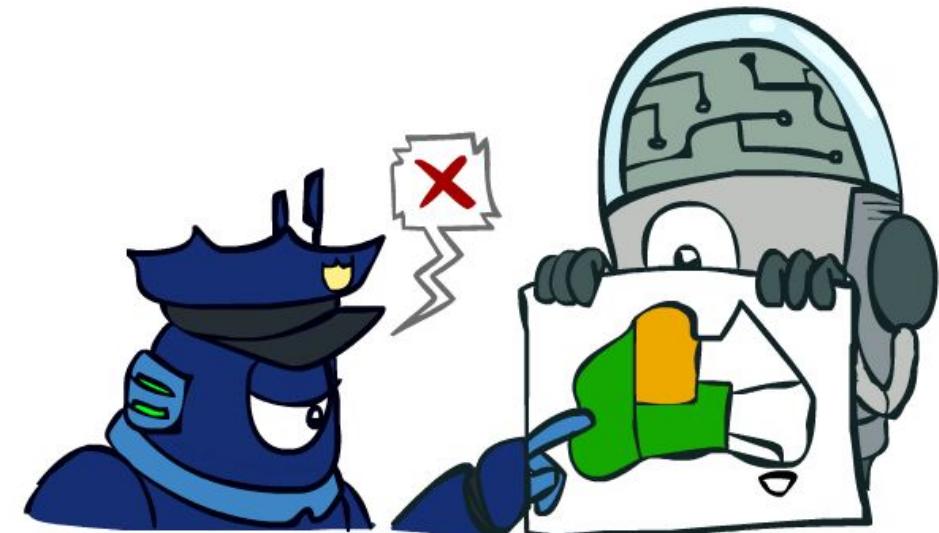
# Backtracking Search

---

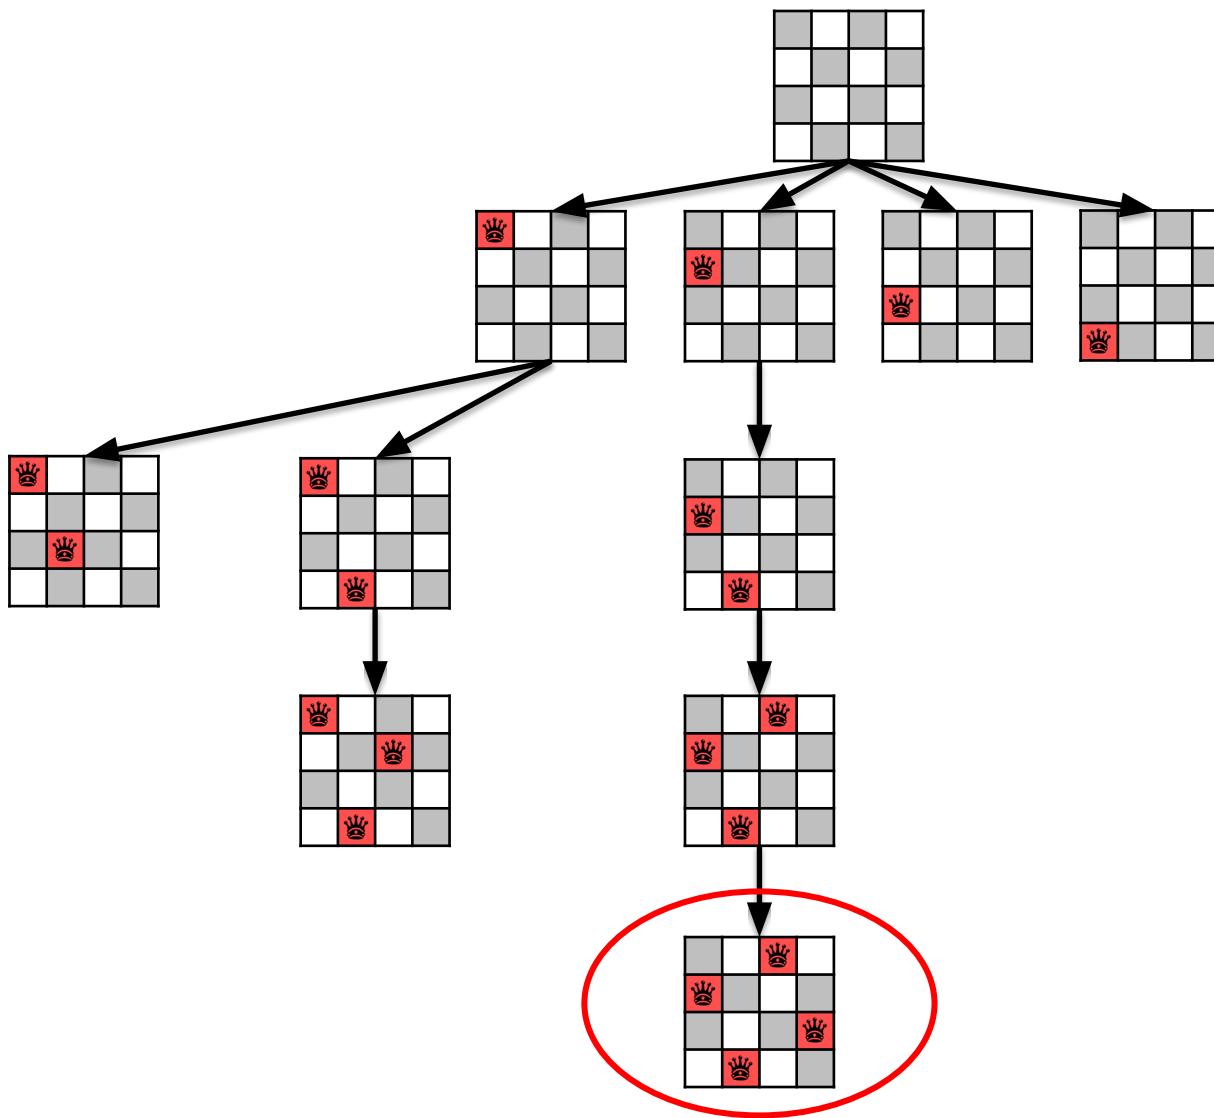


# Backtracking Search

- Backtracking search is the basic uninformed algorithm for solving CSPs
- Idea 1: One variable at a time
  - Only need to consider assignments to a single variable at each step
- Idea 2: Check constraints as you go
  - i.e. consider only values which do not conflict previous assignments
  - Might have to do some computation to check the constraints
  - “Incremental goal test”
- Depth-first search with these two improvements is called *backtracking search* (not the best name)
- Can solve n-queens for  $n \approx 25$

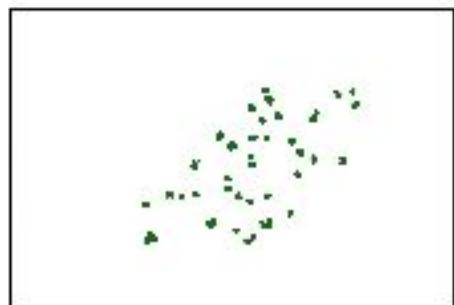


# Backtracking Example



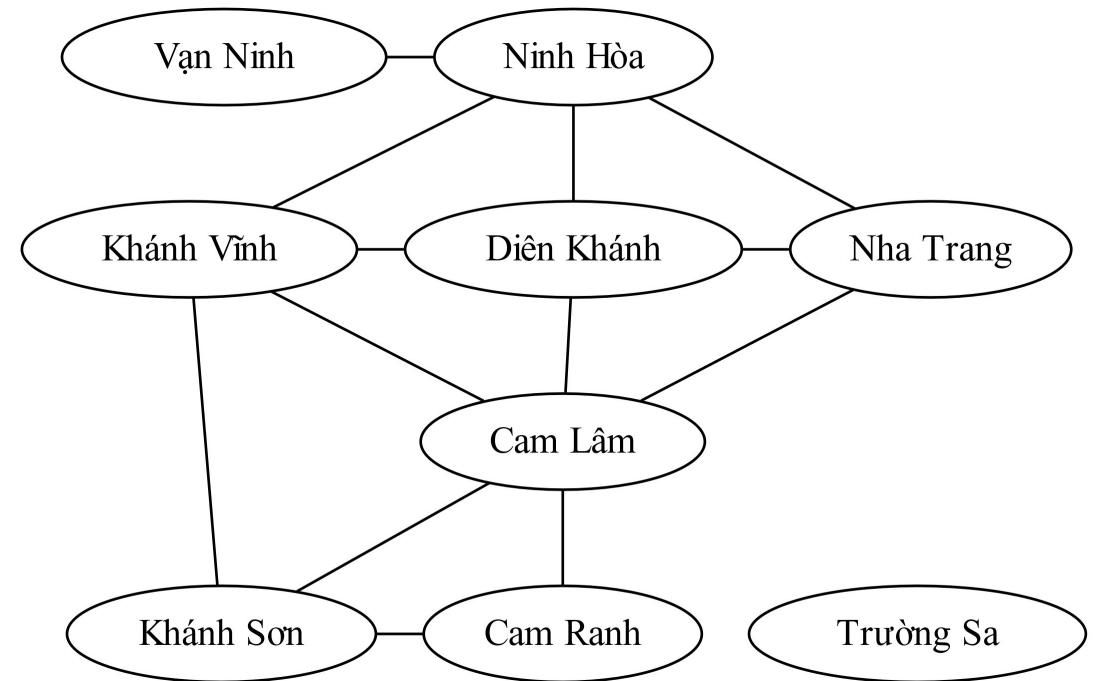
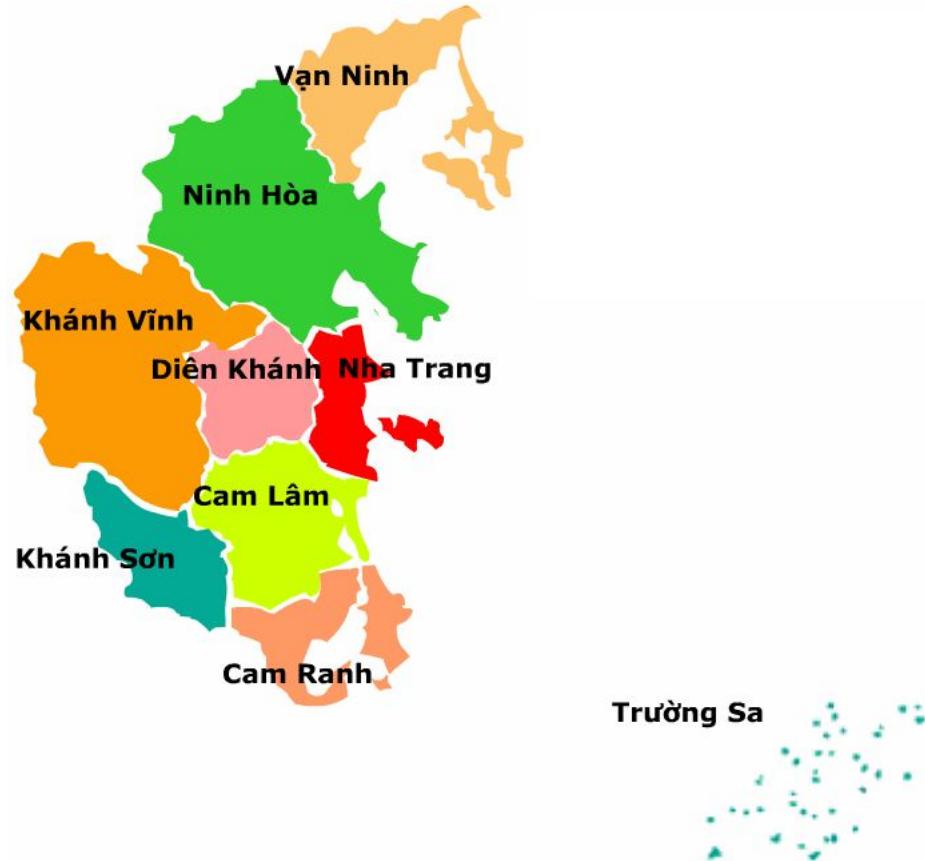
# Map Coloring with Backtracking

---

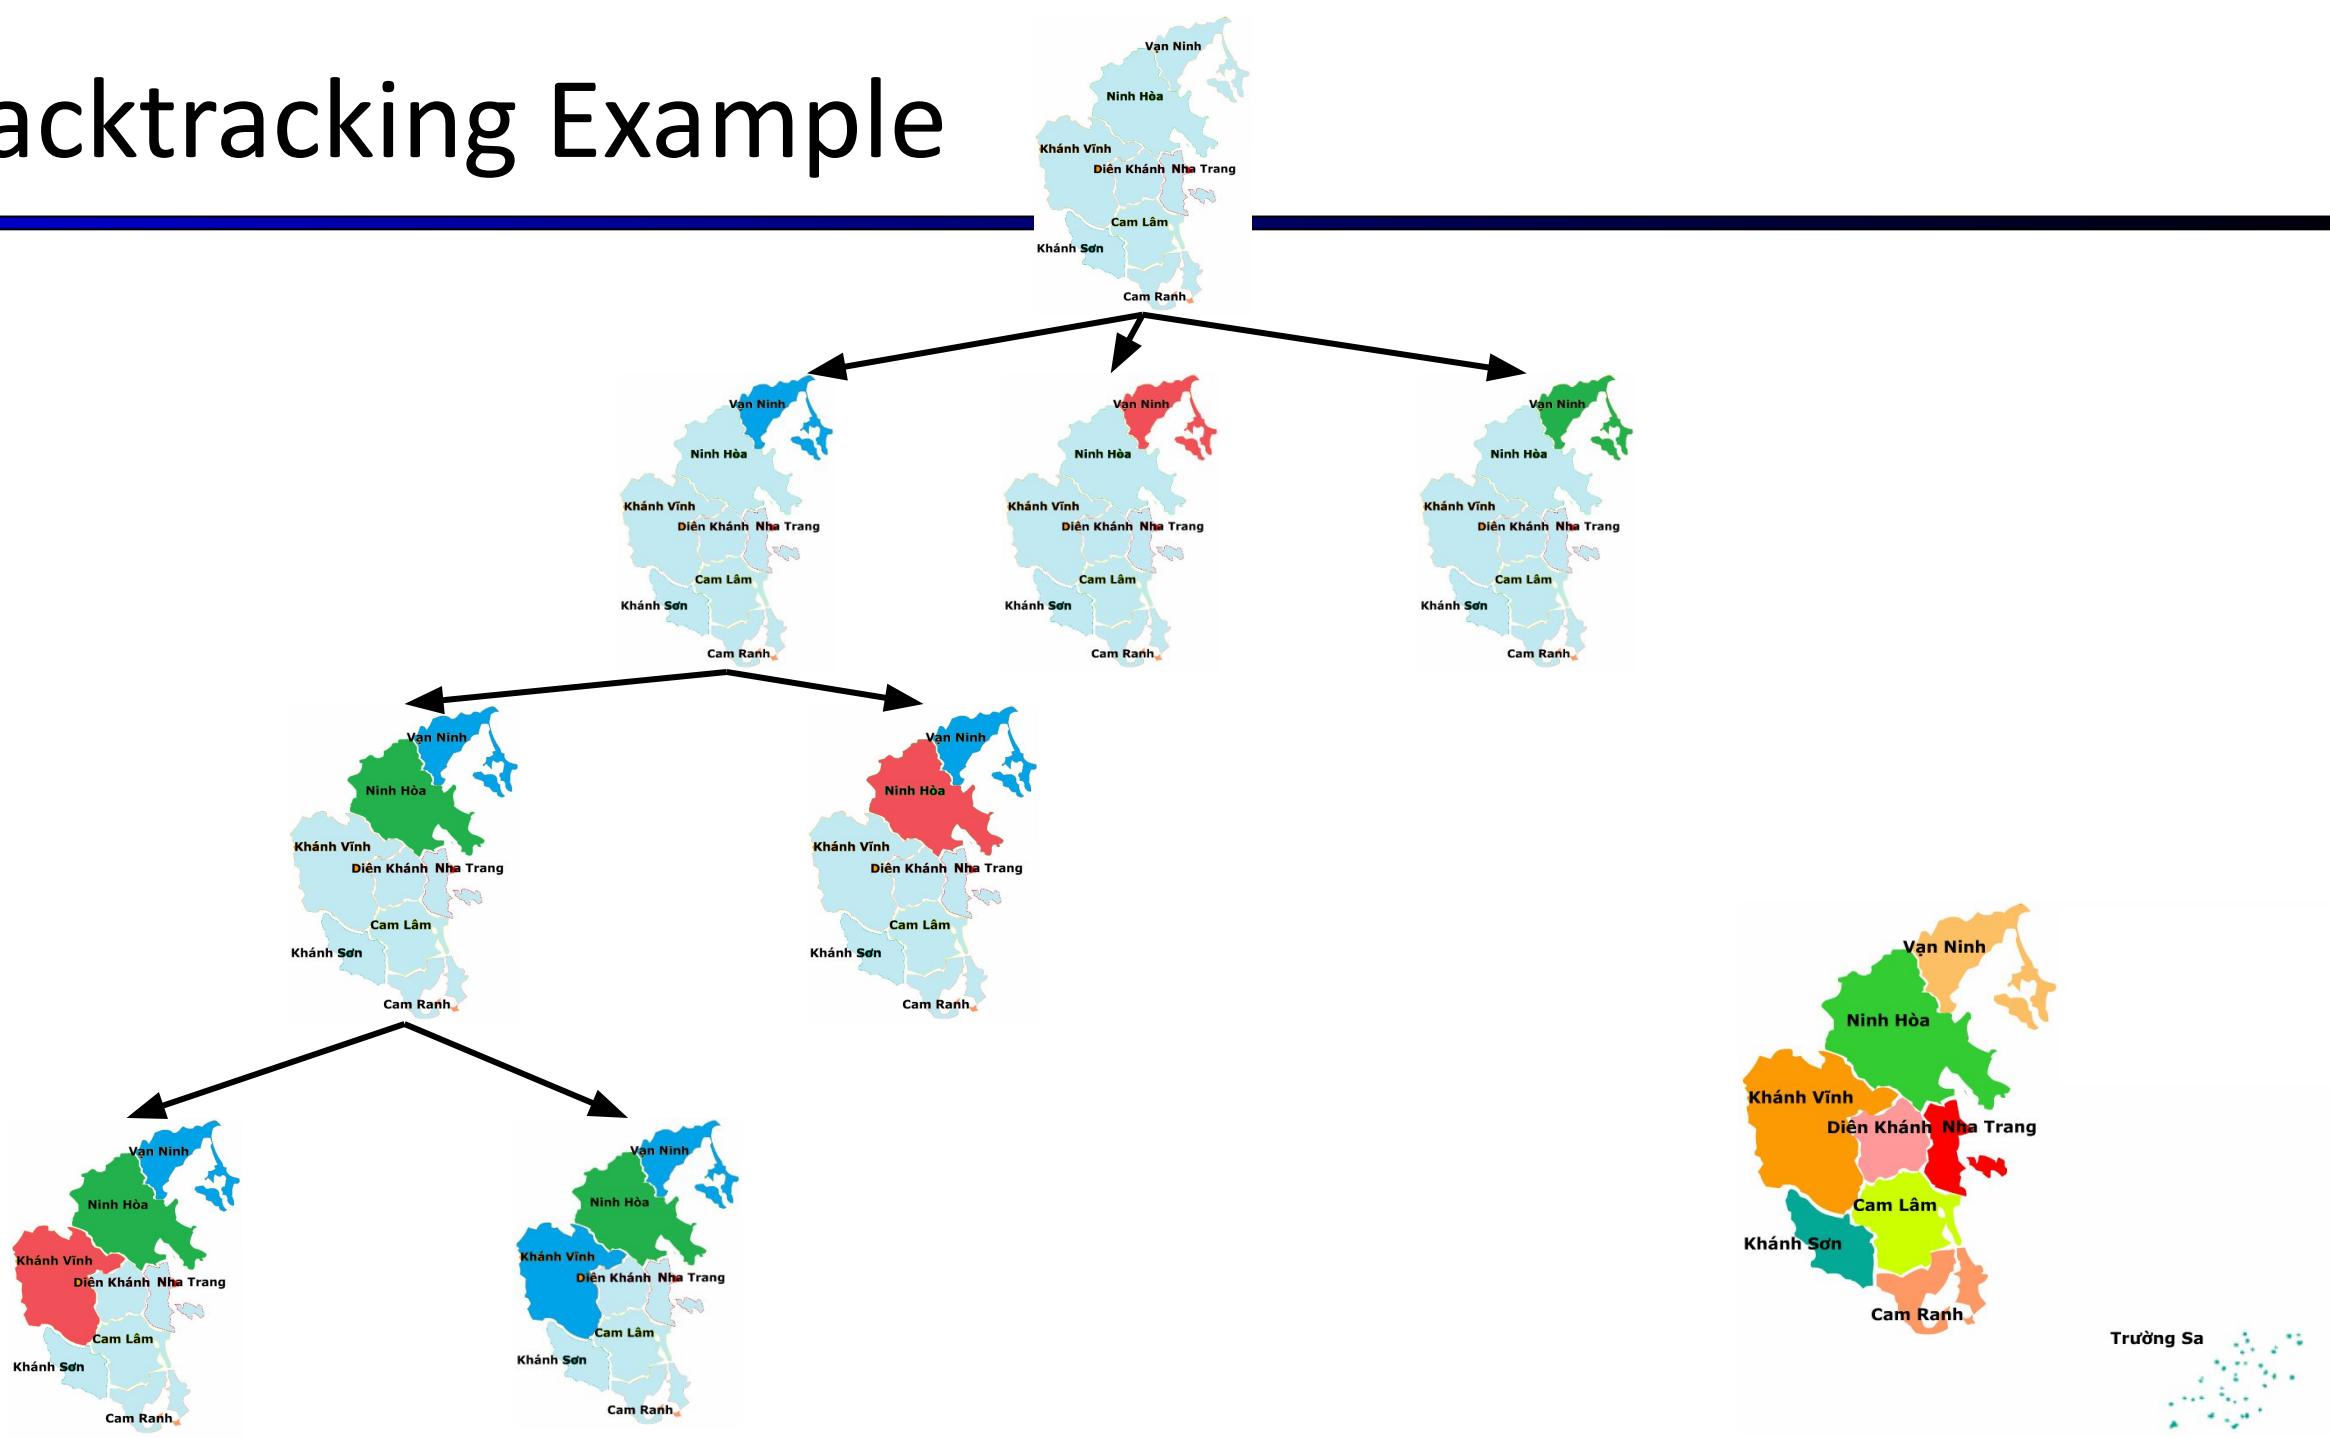


**Where are we  
now?**

# Map Coloring with Backtracking



# Backtracking Example



# Backtracking Search

```
CSP-BACKTRACKING(PartialAssignment a)
```

```
    If a is complete then return a
```

```
    X <- select an unassigned variable
```

```
    D <- select an ordering for the domain of X
```

```
    For each value v in D do
```

```
        If v is consistent with a then
```

```
            Add (X = v) to a
```

```
            result <- CSP-BACKTRACKING(a)
```

```
            If result <> failure then return result
```

```
            Remove (X = v) from a
```

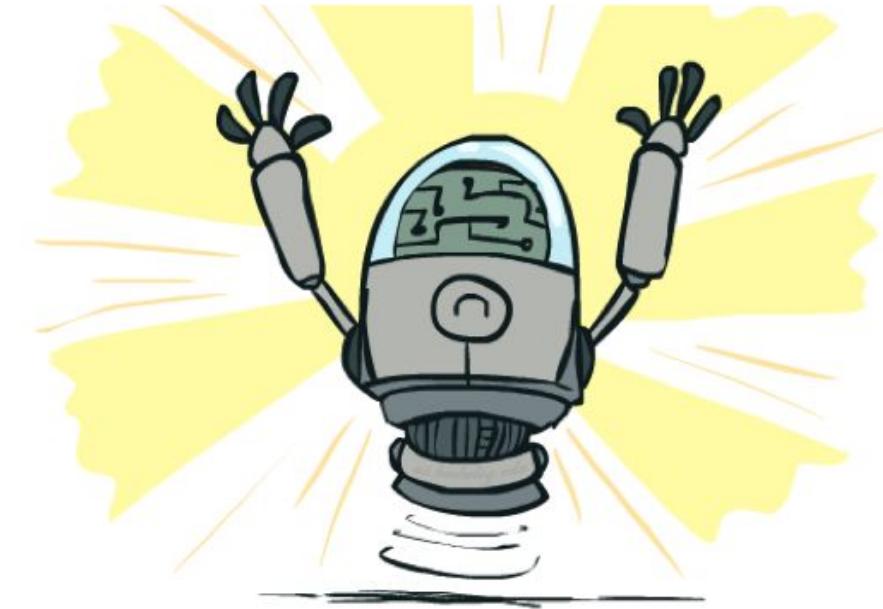
```
    Return failure
```

- Backtracking = DFS + variable-ordering + fail-on-violation
- What are the choice points?

# Improving Backtracking

---

- General-purpose ideas give huge gains in speed
- Filtering: Can we detect inevitable failure early?
- Ordering:
  - Which variable should be assigned next?
  - In what order should its values be tried?
- Structure: Can we exploit the problem structure?



# Filtering

---

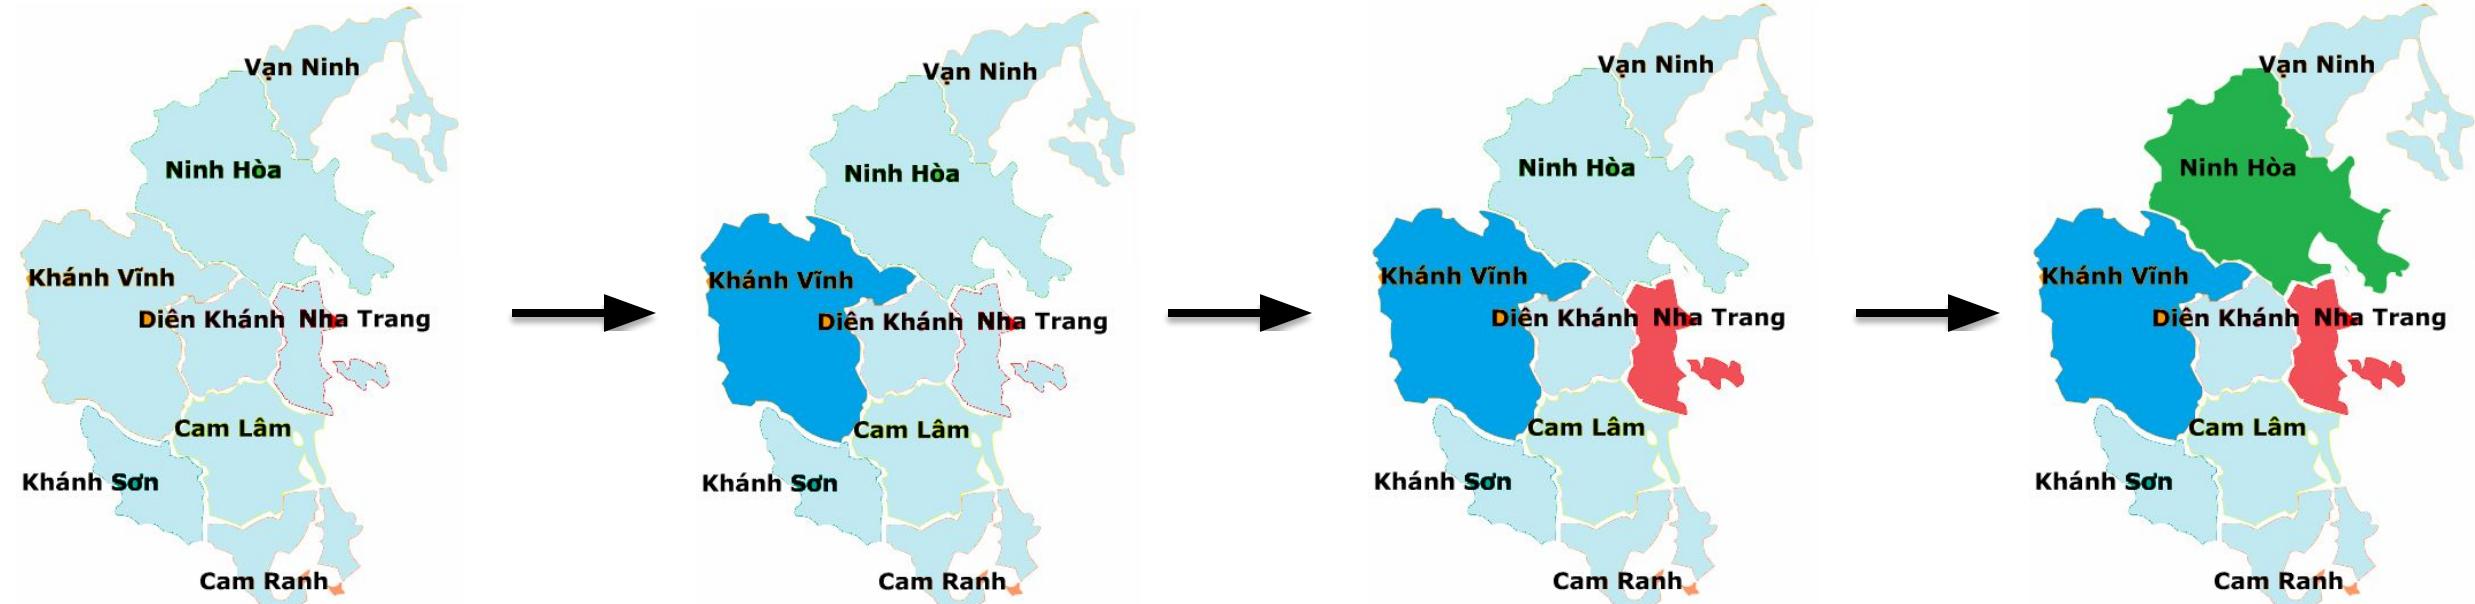


# Filtering: Forward Checking

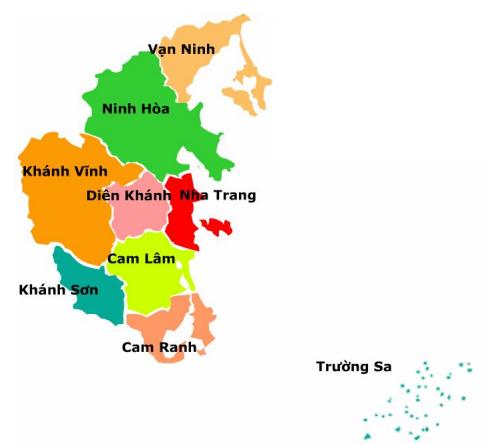
---

- Filtering: Keep track of domains for unassigned variables and cross off bad options
- Forward checking: Cross off values that violate a constraint when added to the existing assignment

# Filtering: Forward Checking



Vạn Ninh	Ninh Hòa	Khánh Vĩnh	Diên Khánh	Nha Trang	Cam Lâm	Khánh Sơn	Cam Ranh	Trường Sa
Red	Green	Blue	Red	Green	Blue	Red	Green	Blue



# Filtering: Constraint Propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:



Vạn Ninh	Ninh Hòa	Khánh Vĩnh	Diên Khánh	Nha Trang	Cam Lâm	Khánh Sơn	Cam Ranh	Trường Sa
Red	Green	Blue	Red	Green	Blue	Red	Green	Blue
Red	Green	Blue	Red	White	Blue	Red	Green	Blue
Red	Green	Blue	White	Green	Blue	Red	White	Blue

- Ninh Hòa and Diên Khánh cannot both be green!
- Why didn't we detect this yet?
- Constraint propagation: reason from constraint to constraint*

# Consistency of A Single Arc

- An arc  $X \rightarrow Y$  is **consistent** iff for *every*  $x$  in the tail there is *some*  $y$  in the head which could be assigned without violating a constraint



Vạn Ninh	Ninh Hòa	Khánh Vĩnh	Diên Khánh	Nha Trang	Cam Lâm	Khánh Sơn	Cam Ranh	Trường Sa
Red	Green	Blue	Blue	Red	Green	Blue	Red	Blue



- Forward checking: Enforcing consistency of arcs pointing to each new assignment

# Consistency of A Single Arc

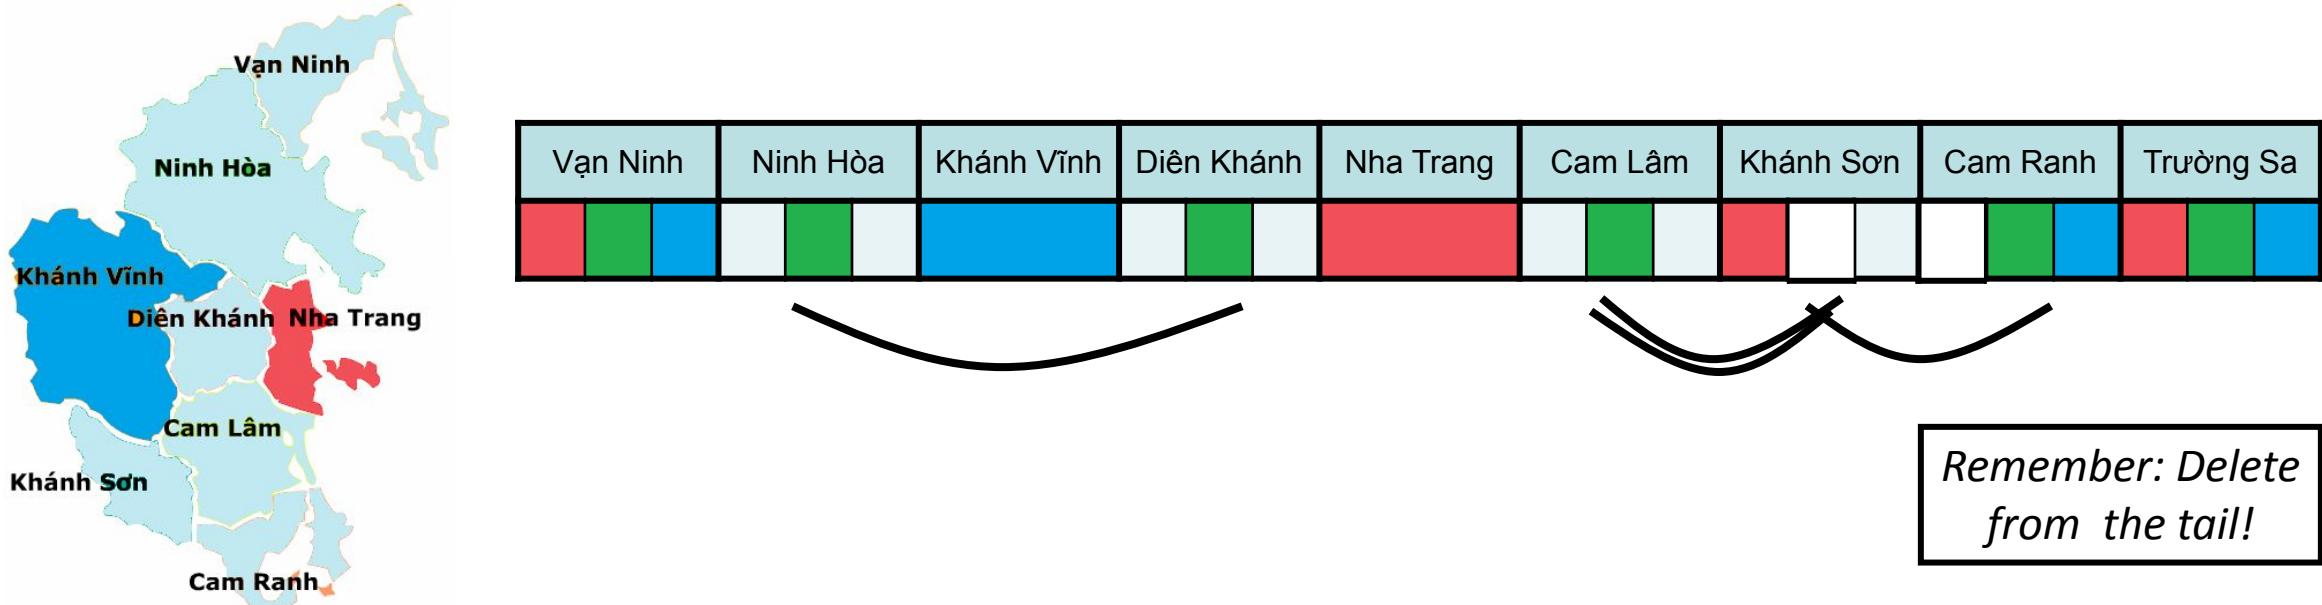
---

- $\text{Revise}(i, j)$  removes values from  $d_i$  without support in  $d_j$ .

```
function Revise( $i, j$ )
    change := false
    for each  $a \in d_i$  do
        if  $\forall_{b \in d_j} \neg c_{ij}(a, b)$  then
            change := true
            remove  $a$  from  $d_i$ 
    return change
```

# Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:



- Important: If  $X$  loses a value, neighbors of  $X$  need to be rechecked!
- Arc consistency detects failure earlier than forward checking
- Can be run as a preprocessor or after each assignment
- What's the downside of enforcing arc consistency?

# Enforcing Arc Consistency in a CSP

---

**procedure** AC-3( $X, D, C$ )

$Q := \{(i, j), (j, i) \mid c_{ij} \in C\}$  // each pair is added twice

**while**  $Q \neq \emptyset$  **do**

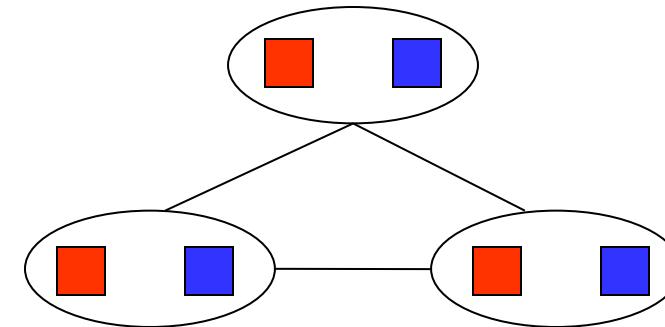
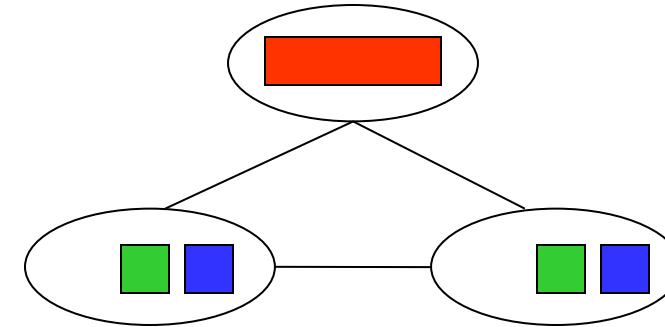
$(i, j) := \text{Fetch}(Q)$  // selects and removes from  $Q$

**if** Revise( $i, j$ ) **then**

$Q := Q \cup \{(k, i) \mid c_{ki} \in C, k \neq j\}$

# Limitations of Arc Consistency

- After enforcing arc consistency:
  - Can have one solution left
  - Can have multiple solutions left
  - Can have no solutions left (and not know it)
- Arc consistency still runs inside a backtracking search!



*What went wrong here?*

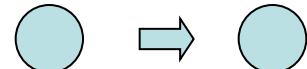
# K-Consistency

- Increasing degrees of consistency

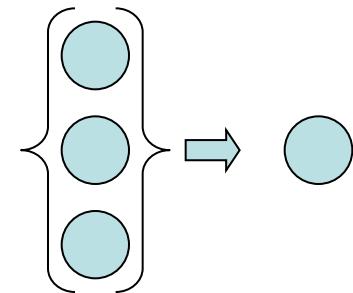
- 1-Consistency (Node Consistency): Each single node's domain has a value which meets that node's unary constraints



- 2-Consistency (Arc Consistency): For each pair of nodes, any consistent assignment to one can be extended to the other

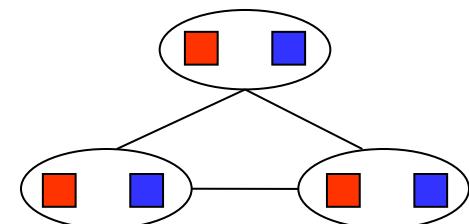


- K-Consistency: For each k nodes, any consistent assignment to k-1 can be extended to the k<sup>th</sup> node.



- Higher k more expensive to compute

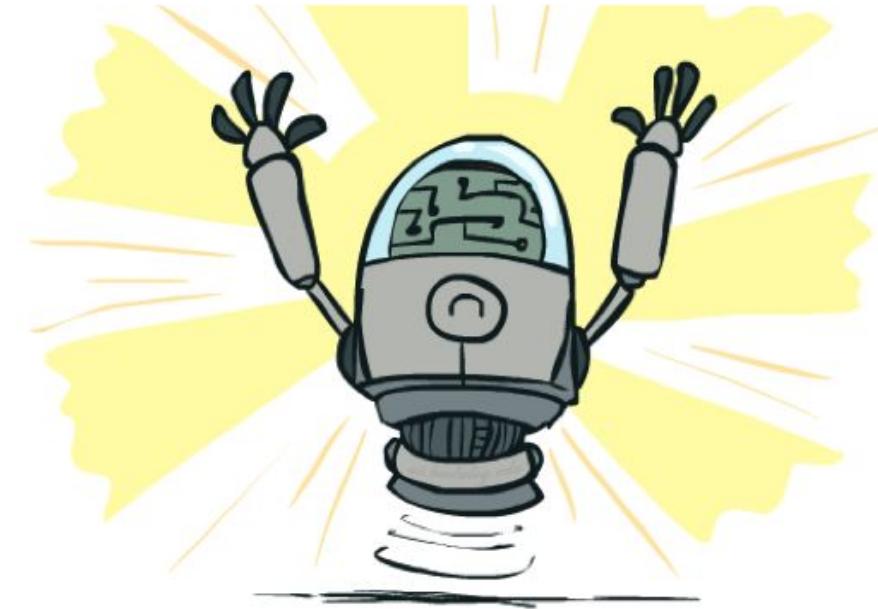
- (You need to know the k=2 case: arc consistency)



# Improving Backtracking

---

- General-purpose ideas give huge gains in speed
- *Filtering: Can we detect inevitable failure early?*
- **Ordering:**
  - Which variable should be assigned next?
  - In what order should its values be tried?
- Structure: Can we exploit the problem structure?



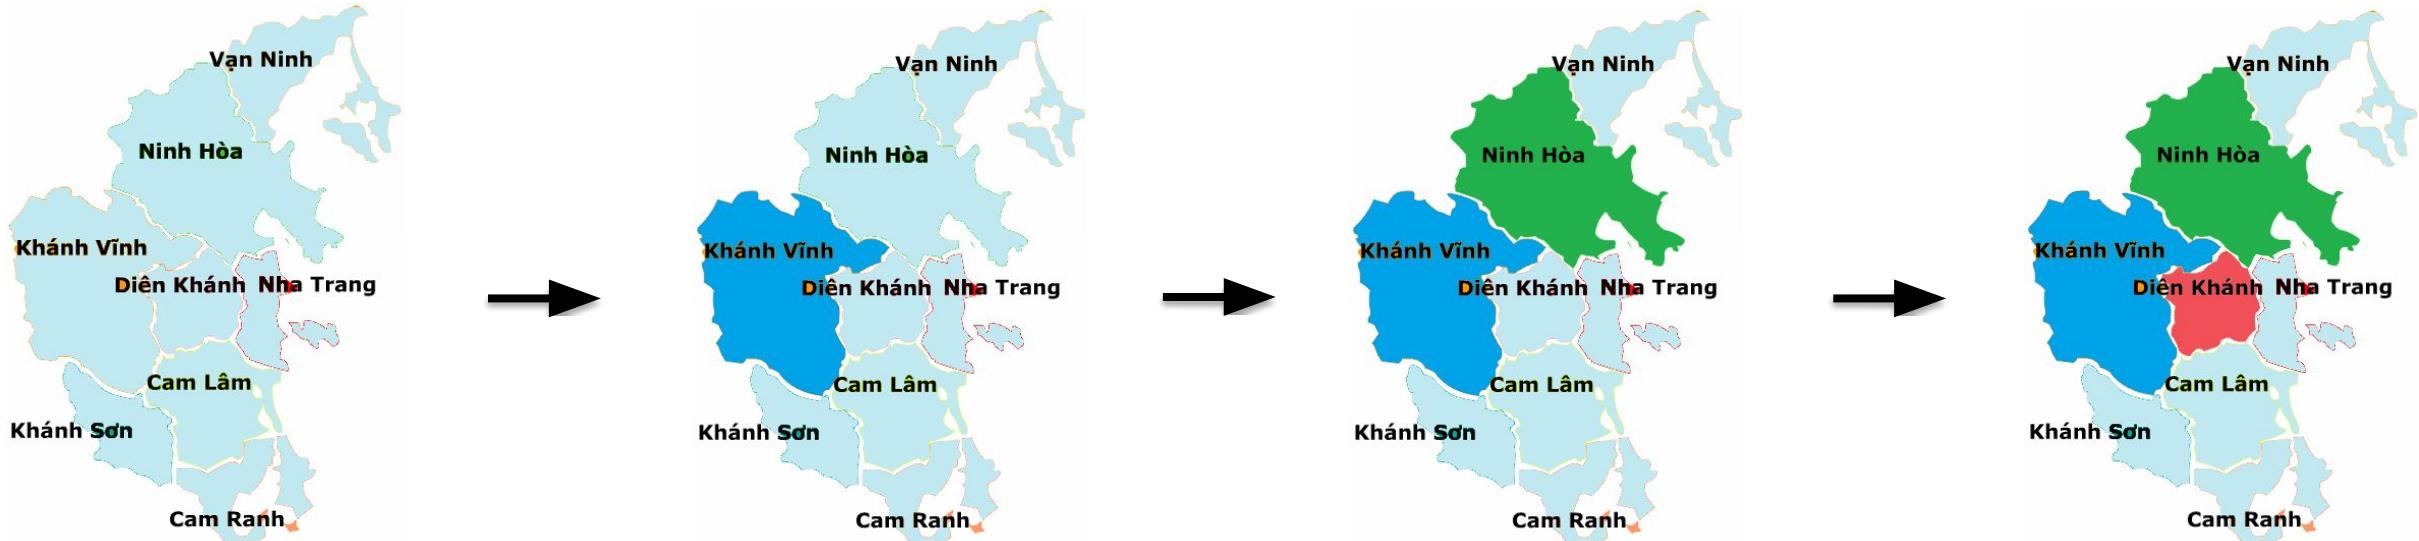
# Ordering

---



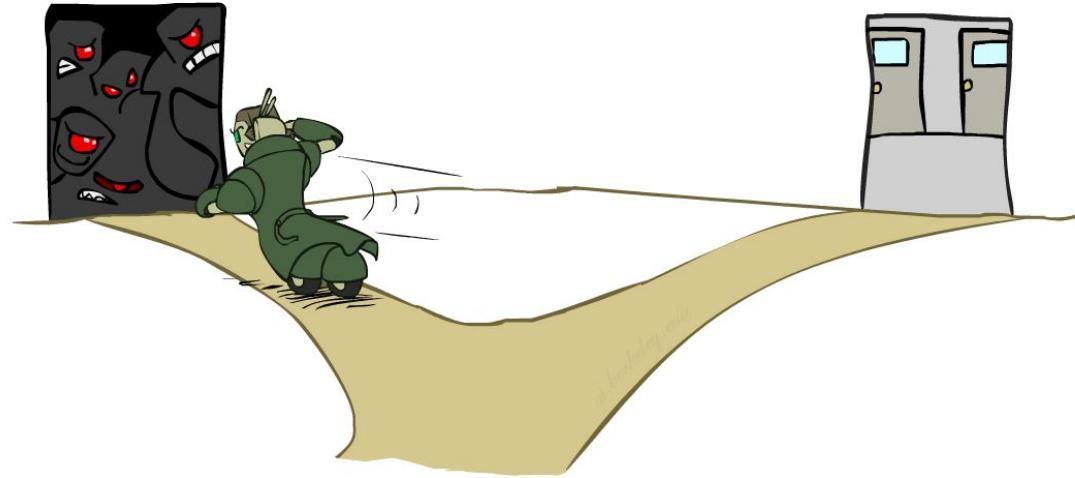
# Ordering: Minimum Remaining Values

- Variable Ordering: Minimum remaining values (MRV):
  - Choose the variable with the fewest legal left values in its domain



# Ordering: Minimum Remaining Values

- Variable Ordering: Minimum remaining values (MRV):
  - Choose the variable with the fewest legal left values in its domain



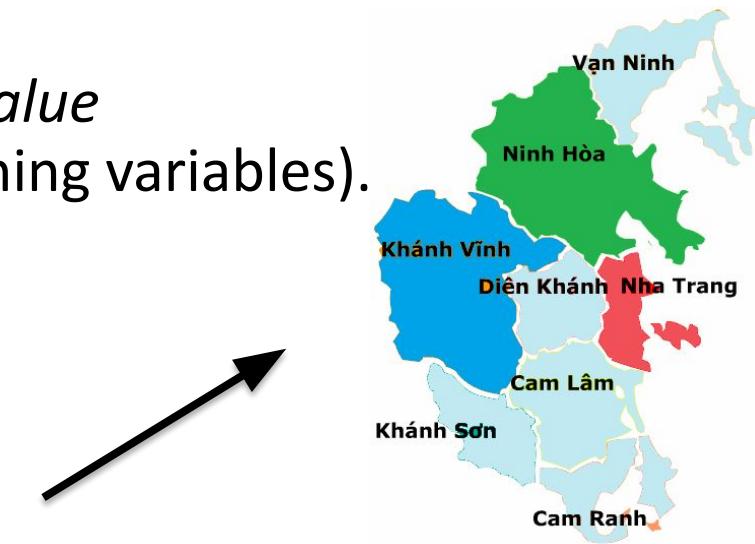
- Why min rather than max?
- Also called “most constrained variable”
- “Fail-fast” ordering

# Ordering: Least Constraining Value

## Value Ordering: Least Constraining Value

- Given a choice of variable, choose the *least constraining value* (i.e., the one that rules out the fewest values in the remaining variables).

Vạn Ninh	Ninh Hòa	Khánh Vĩnh	Diên Khánh	Nha Trang	Cam Lâm	Khánh Sơn	Cam Ranh	Trường Sa
Red	Light Blue	Blue	Red	Light Blue	Red	Green	Light Blue	Red

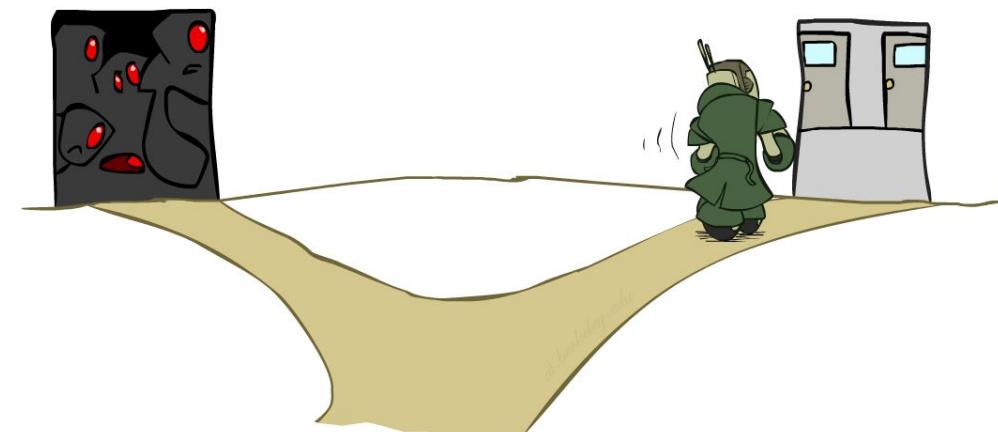
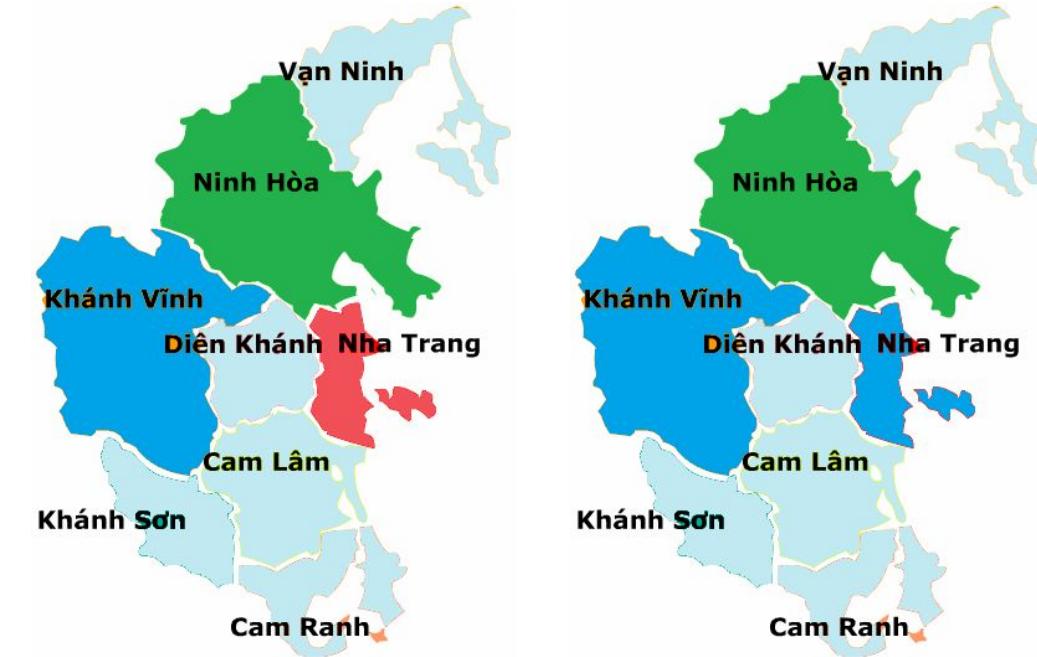


?



# Ordering: Least Constraining Value

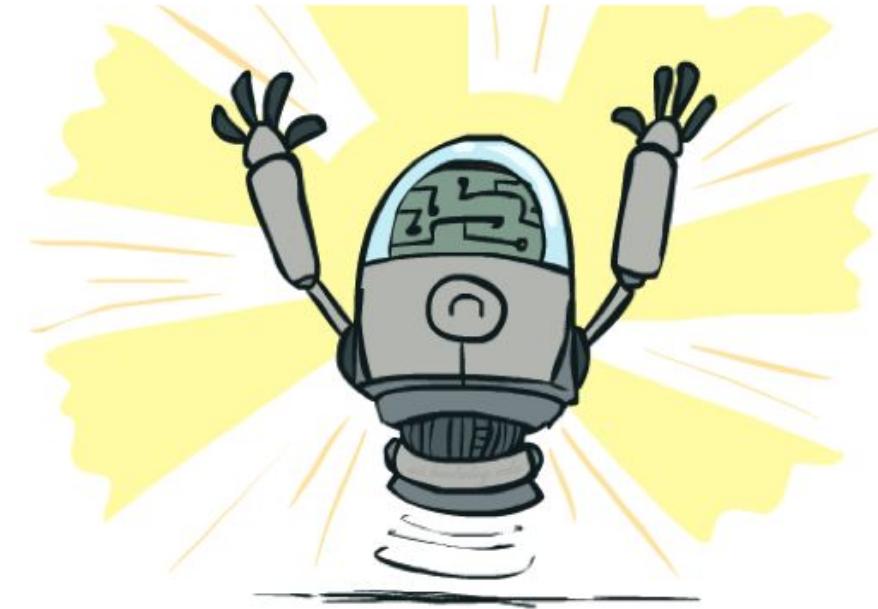
- **Value Ordering: Least Constraining Value**
  - Given a choice of variable, choose the *least constraining value*
  - I.e., the one that rules out the fewest values in the remaining variables
  - **Note that it may take some computation to determine this! (E.g., rerunning filtering)**
- Why least rather than most?
- Combining these ordering ideas makes 1000 queens feasible



# Improving Backtracking

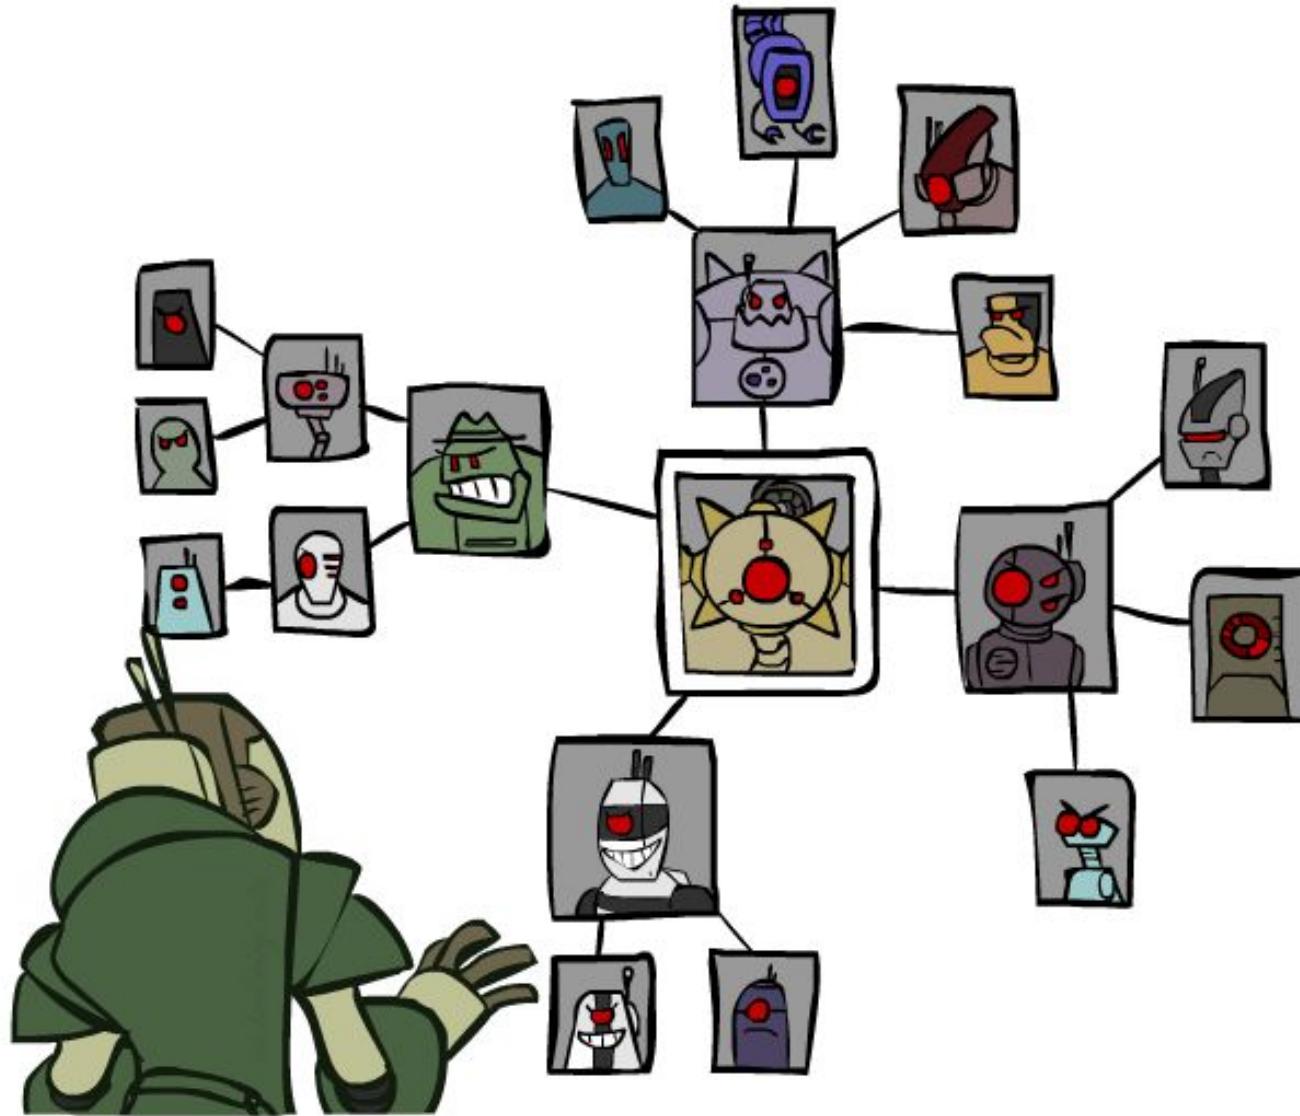
---

- General-purpose ideas give huge gains in speed
- *Filtering: Can we detect inevitable failure early?*
- *Ordering:*
  - *Which variable should be assigned next?*
  - *In what order should its values be tried?*
- **Structure: Can we exploit the problem structure?**



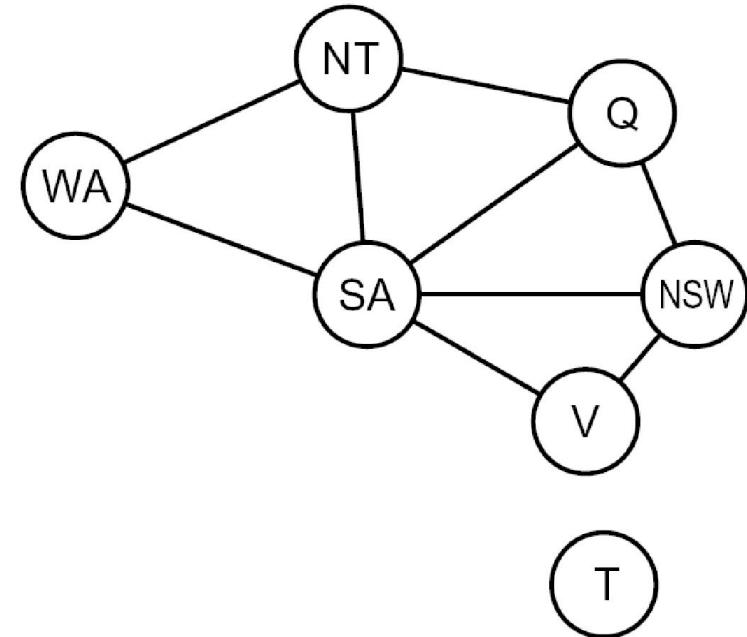
# Structure

---



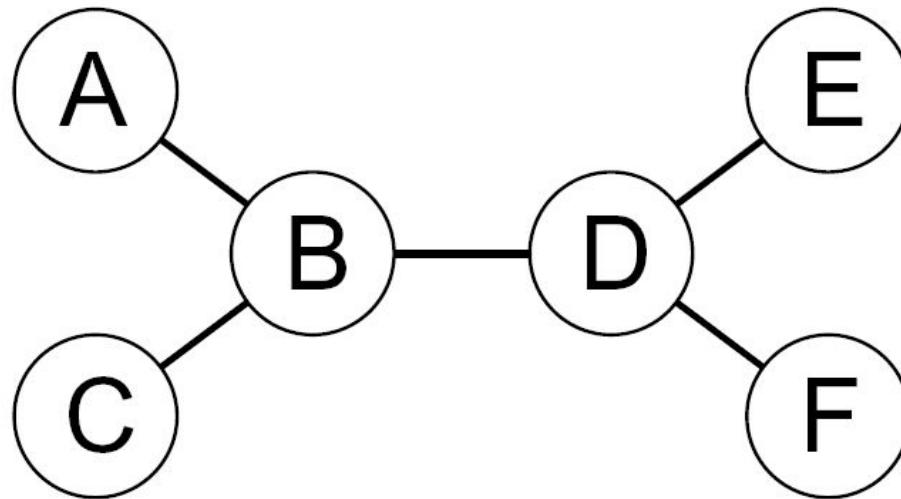
# Problem Structure

- Extreme case: independent subproblems
  - Example: Tasmania and mainland do not interact
- Independent subproblems are identifiable as connected components of constraint graph
- Suppose a graph of  $n$  variables can be broken into subproblems of only  $c$  variables:
  - Worst-case solution cost is  $O((n/c)(d^c))$ , linear in  $n$
  - E.g.,  $n = 80$ ,  $d = 2$ ,  $c = 20$
  - $2^{80} = 4$  billion years at 10 million nodes/sec
  - $(4)(2^{20}) = 0.4$  seconds at 10 million nodes/sec



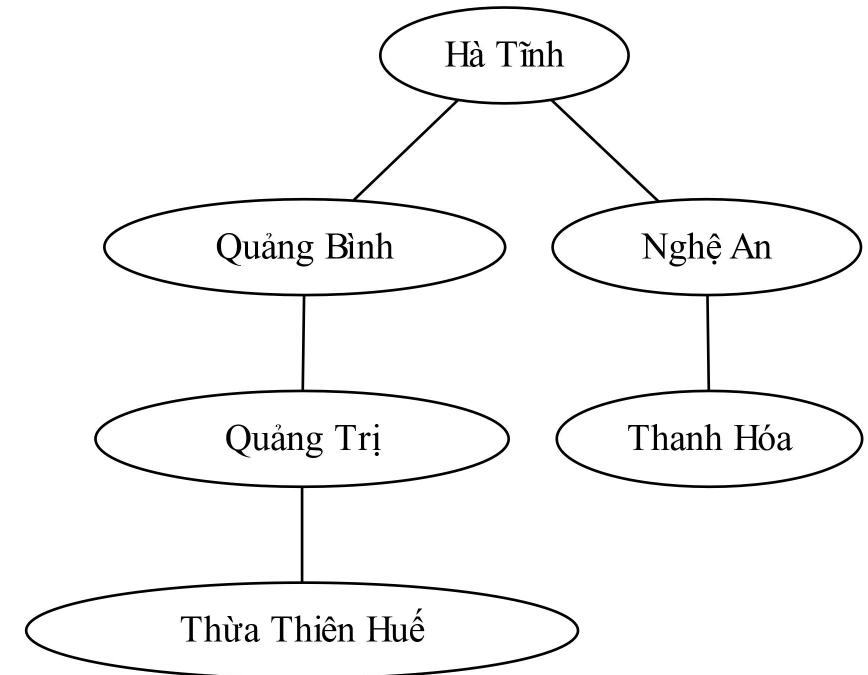
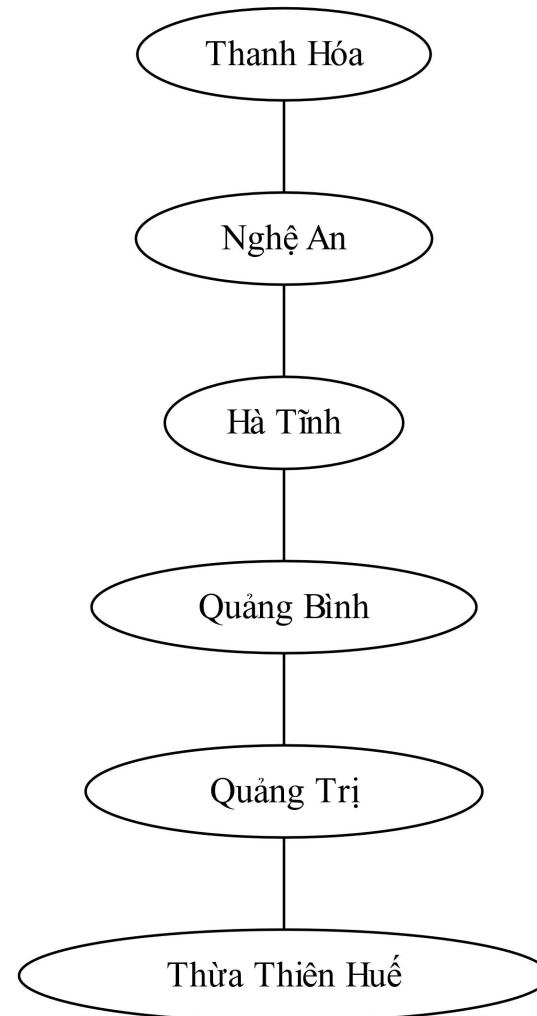
# Tree-Structured CSPs

---



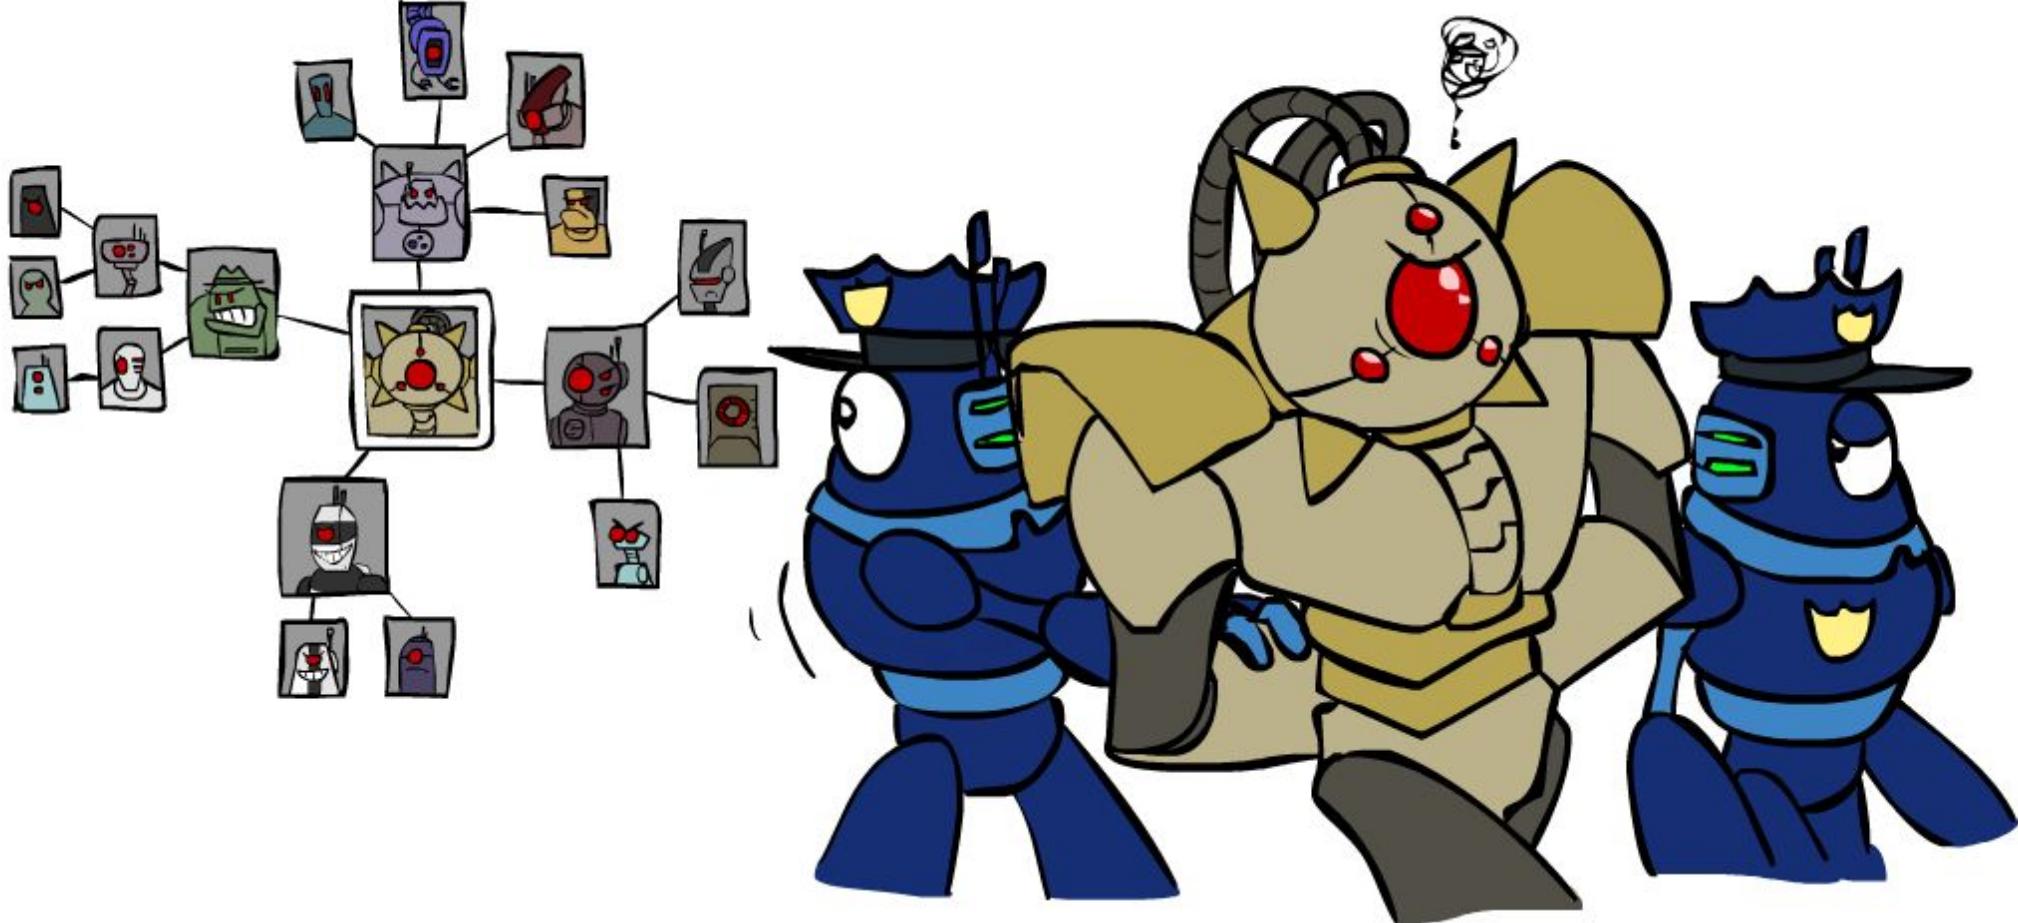
- **Theorem:** if the constraint graph has no loops, the CSP can be solved in  $O(n d^2)$  time
  - Compare to general CSPs, where worst-case time is  $O(d^n)$

# Tree-Structured CSPs

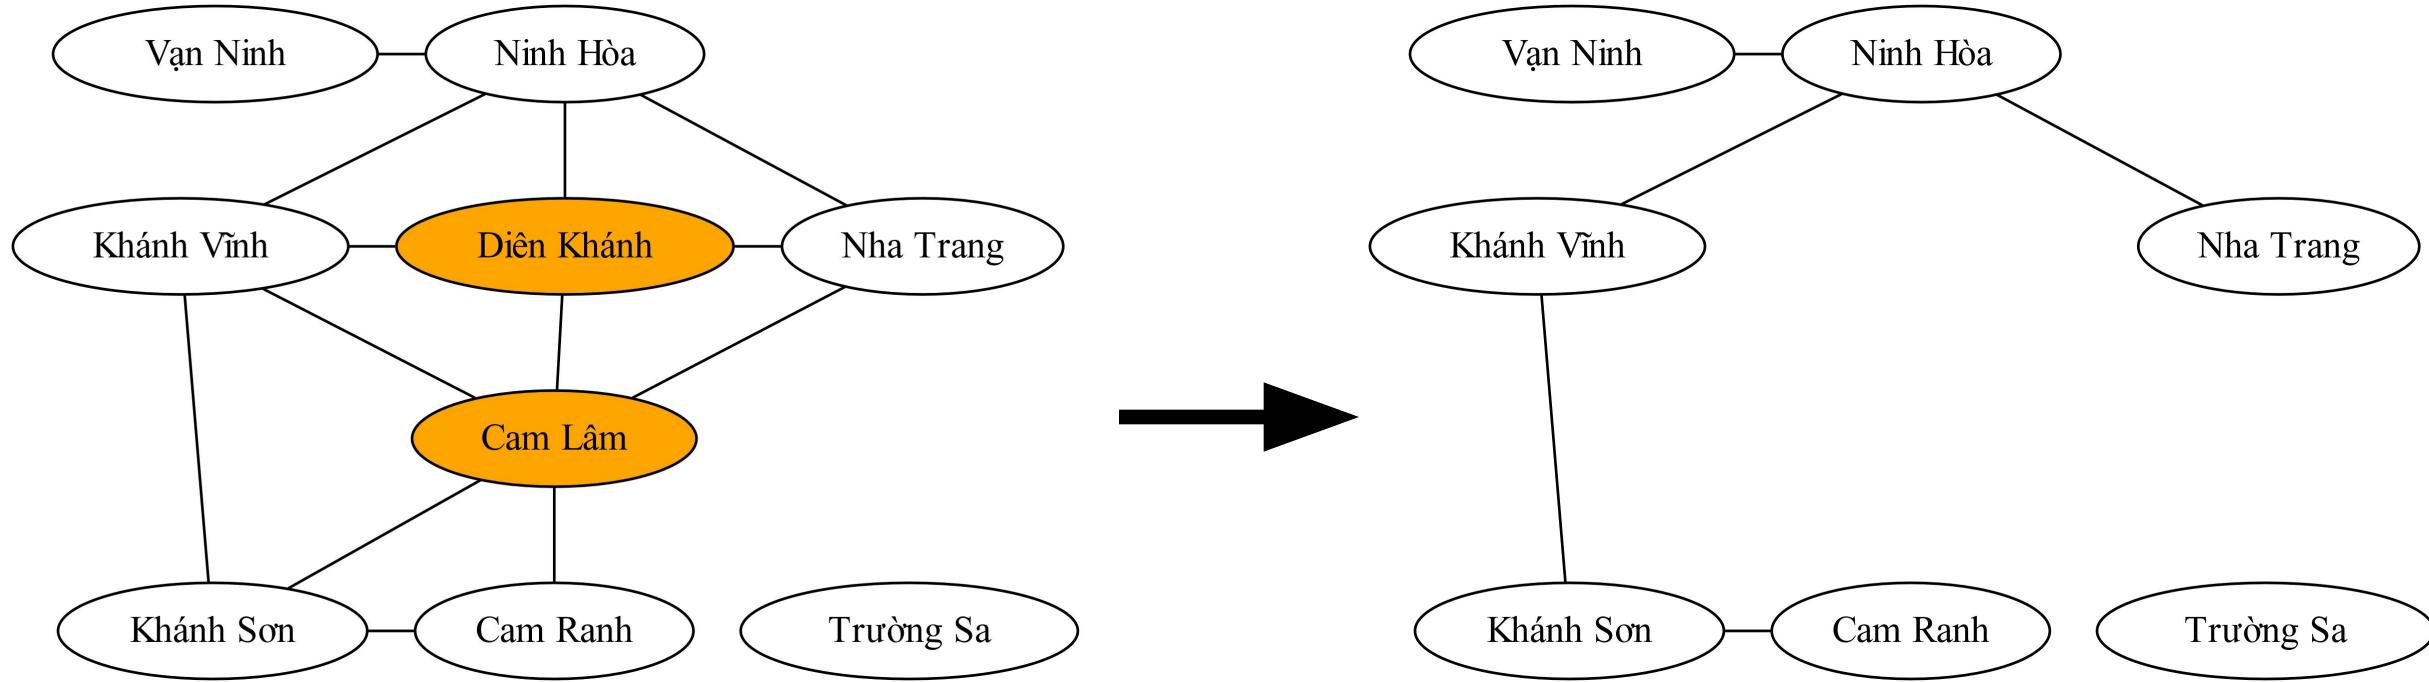


# Improving Structure

---



# Nearly Tree-Structured CSPs



- Conditioning: instantiate a variable, prune its neighbors' domains
- Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree

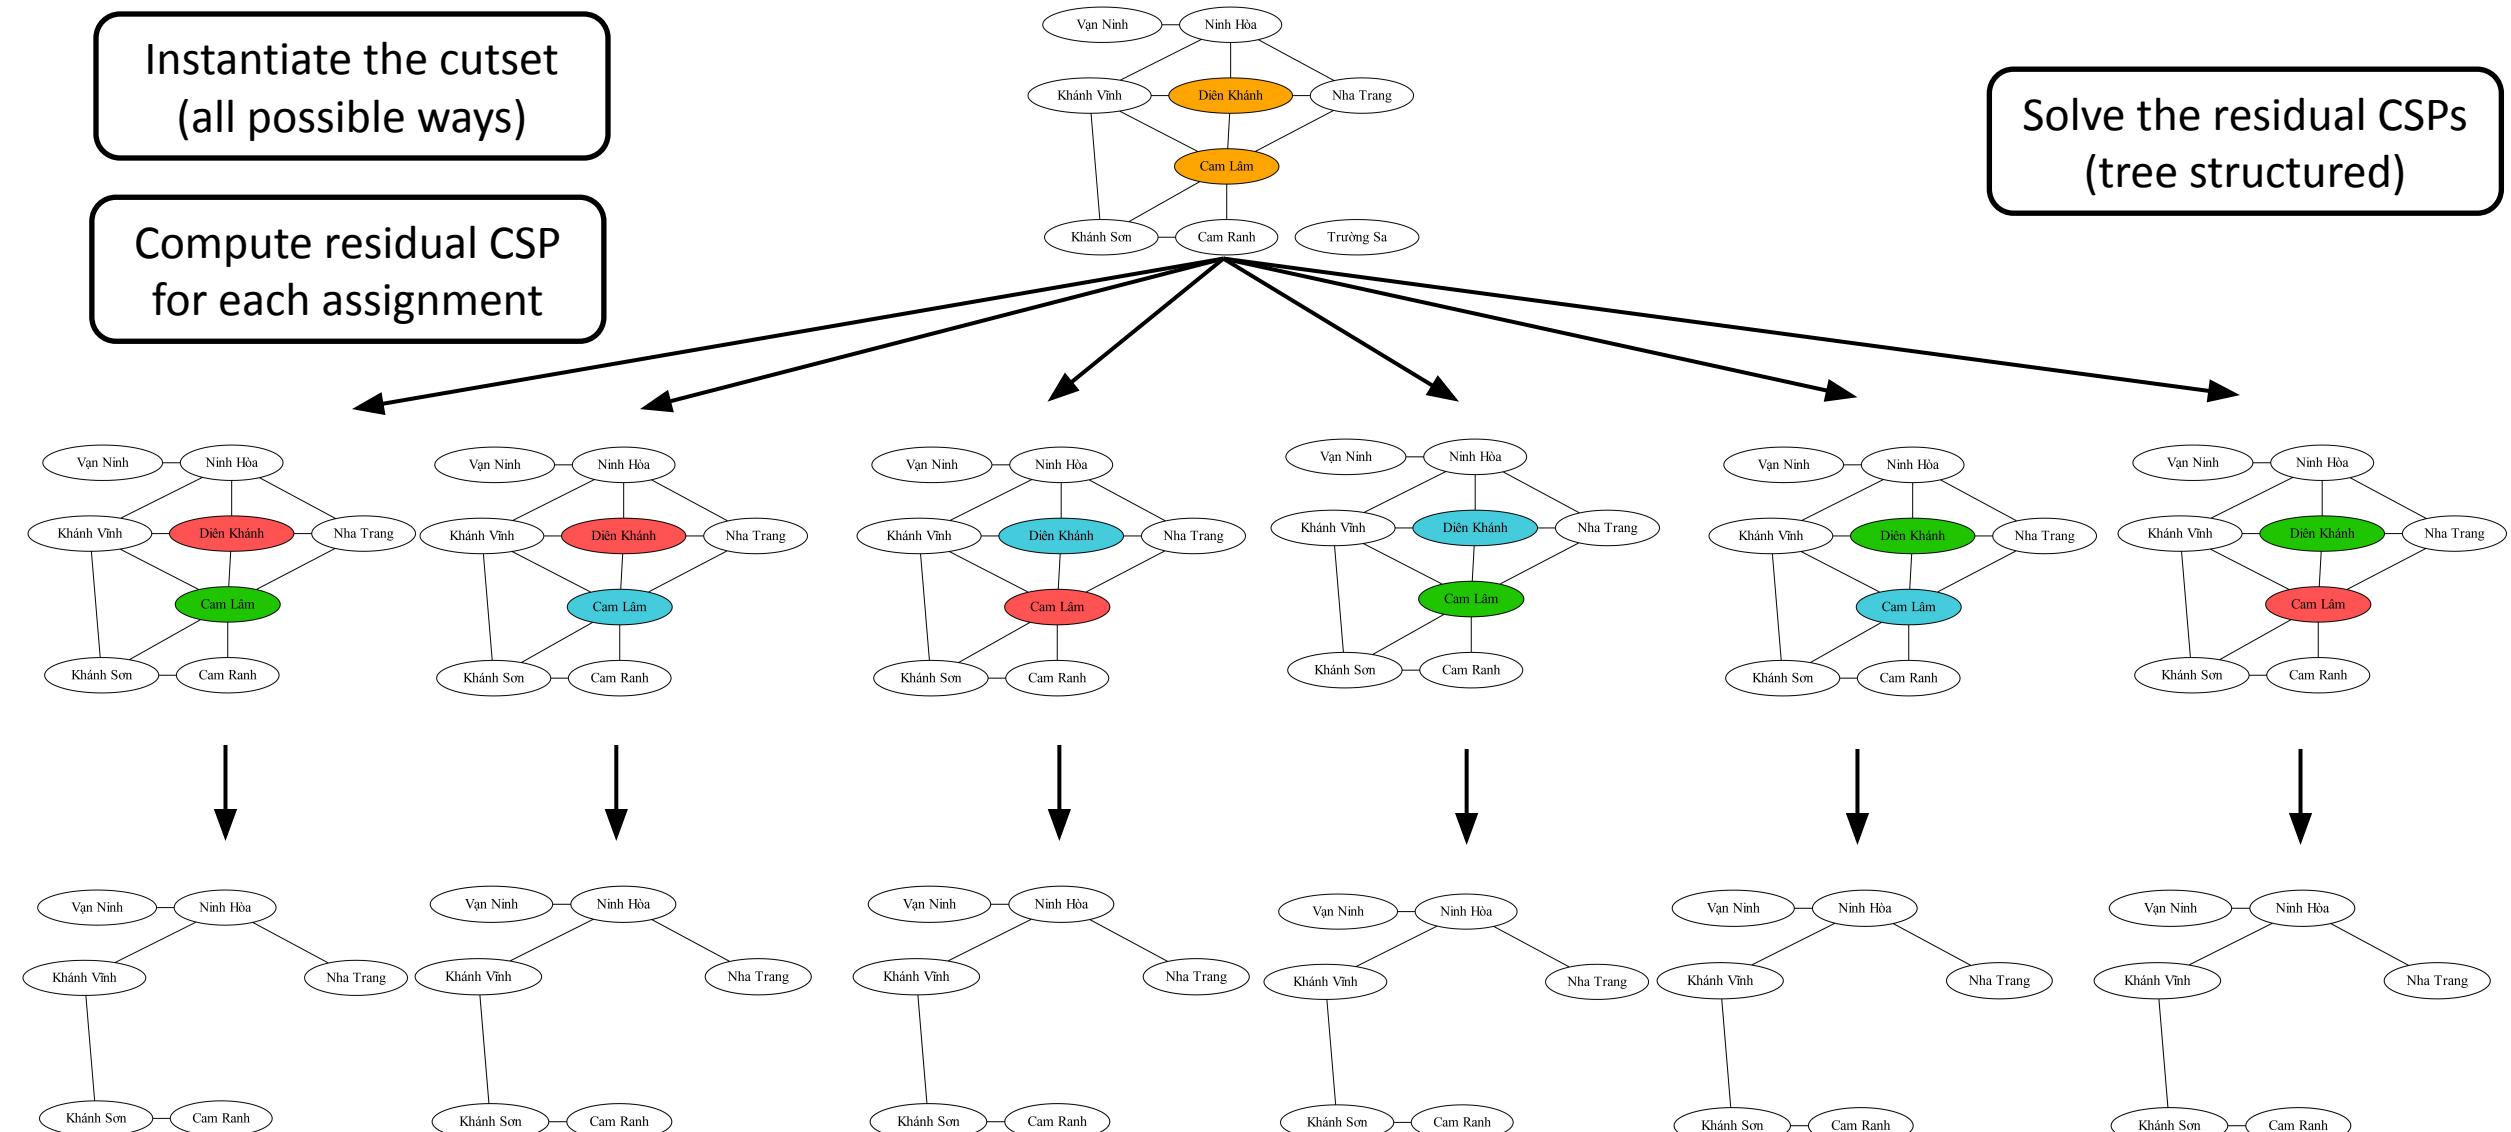
Choose a cutset

# Cutset Conditioning

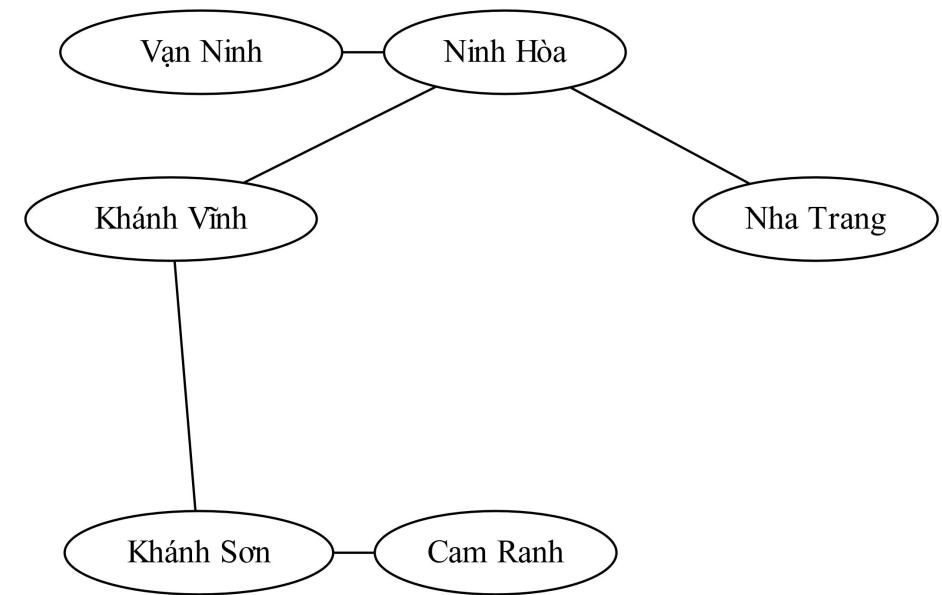
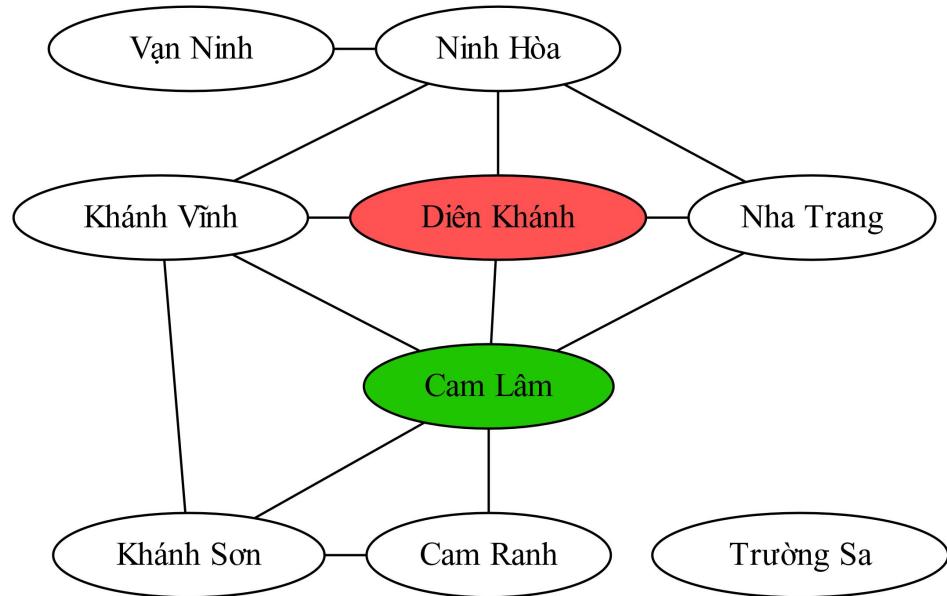
Instantiate the cutset  
(all possible ways)

Compute residual CSP  
for each assignment

Solve the residual CSPs  
(tree structured)



# Cutset Conditioning

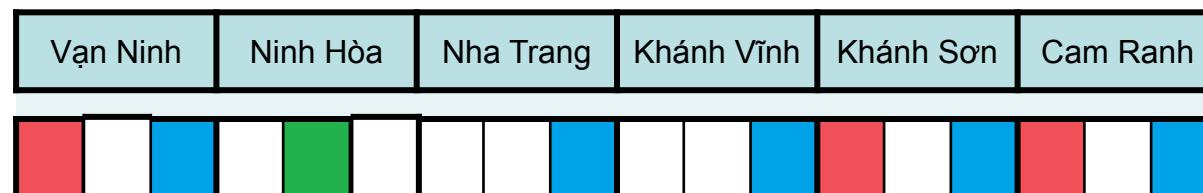
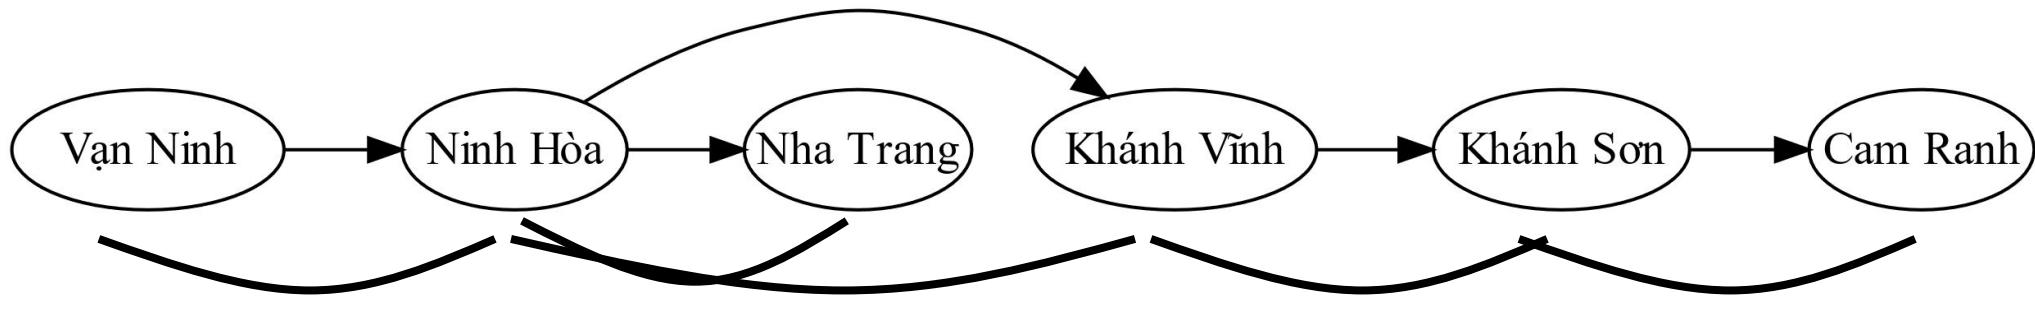


Vạn Ninh	Ninh Hòa	Khánh Vĩnh	Diên Khánh	Nha Trang	Cam Lâm	Khánh Sơn	Cam Ranh	Trường Sa
Red	Green	Blue	Red	Green	Blue	Red	Green	Blue

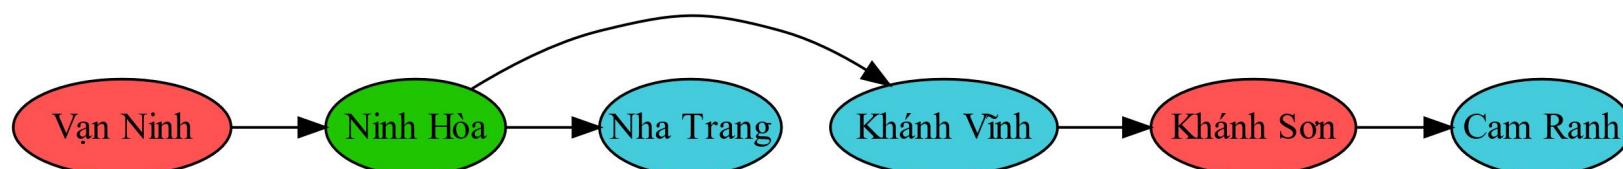
# Tree-Structured CSP

- Algorithm for tree-structured CSPs:

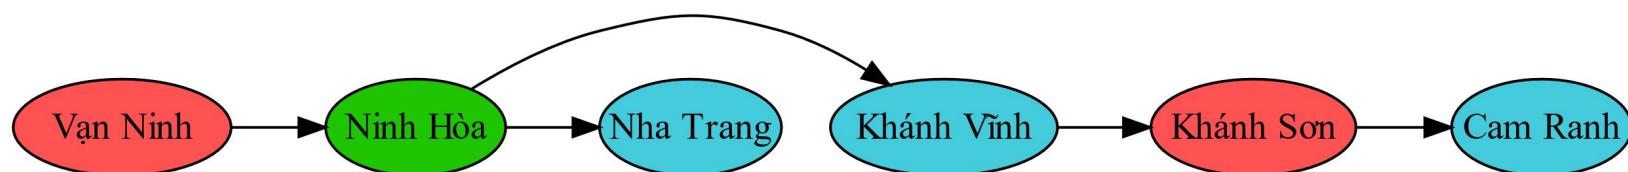
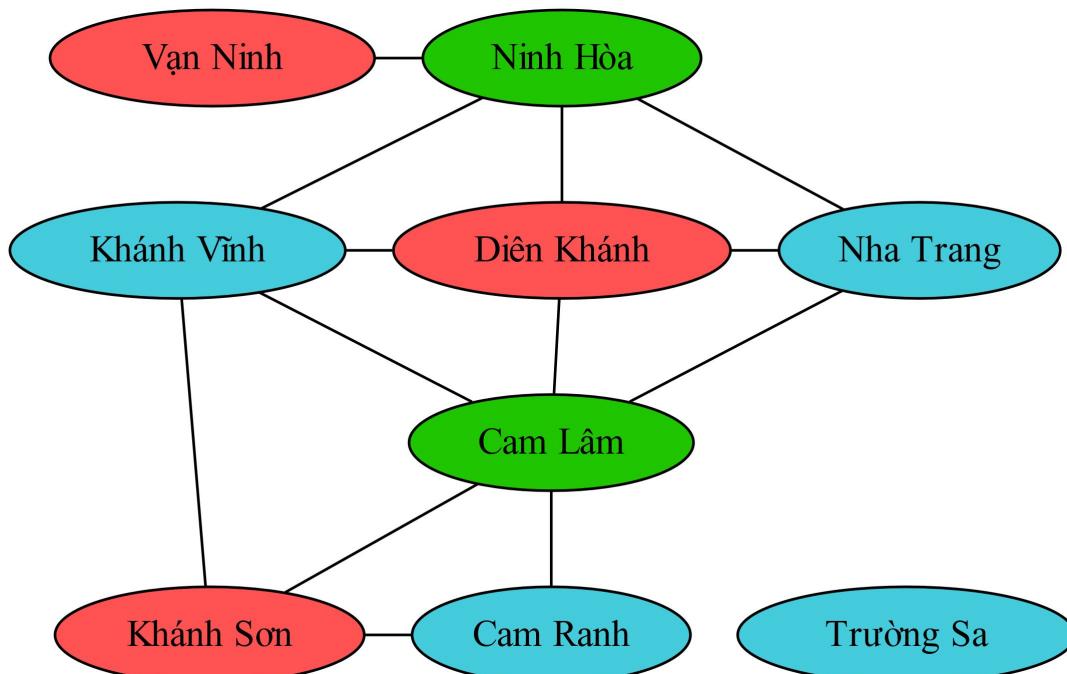
- Order: Choose a root variable, order variables so that parents precede children



- Remove backward: For  $i = n : 2$ , apply Remove-Inconsistency( $\text{Parent}(X_i), X_i$ )
- Assign forward: For  $i = 1 : n$ , assign  $X_i$  consistently with  $\text{Parent}(X_i)$

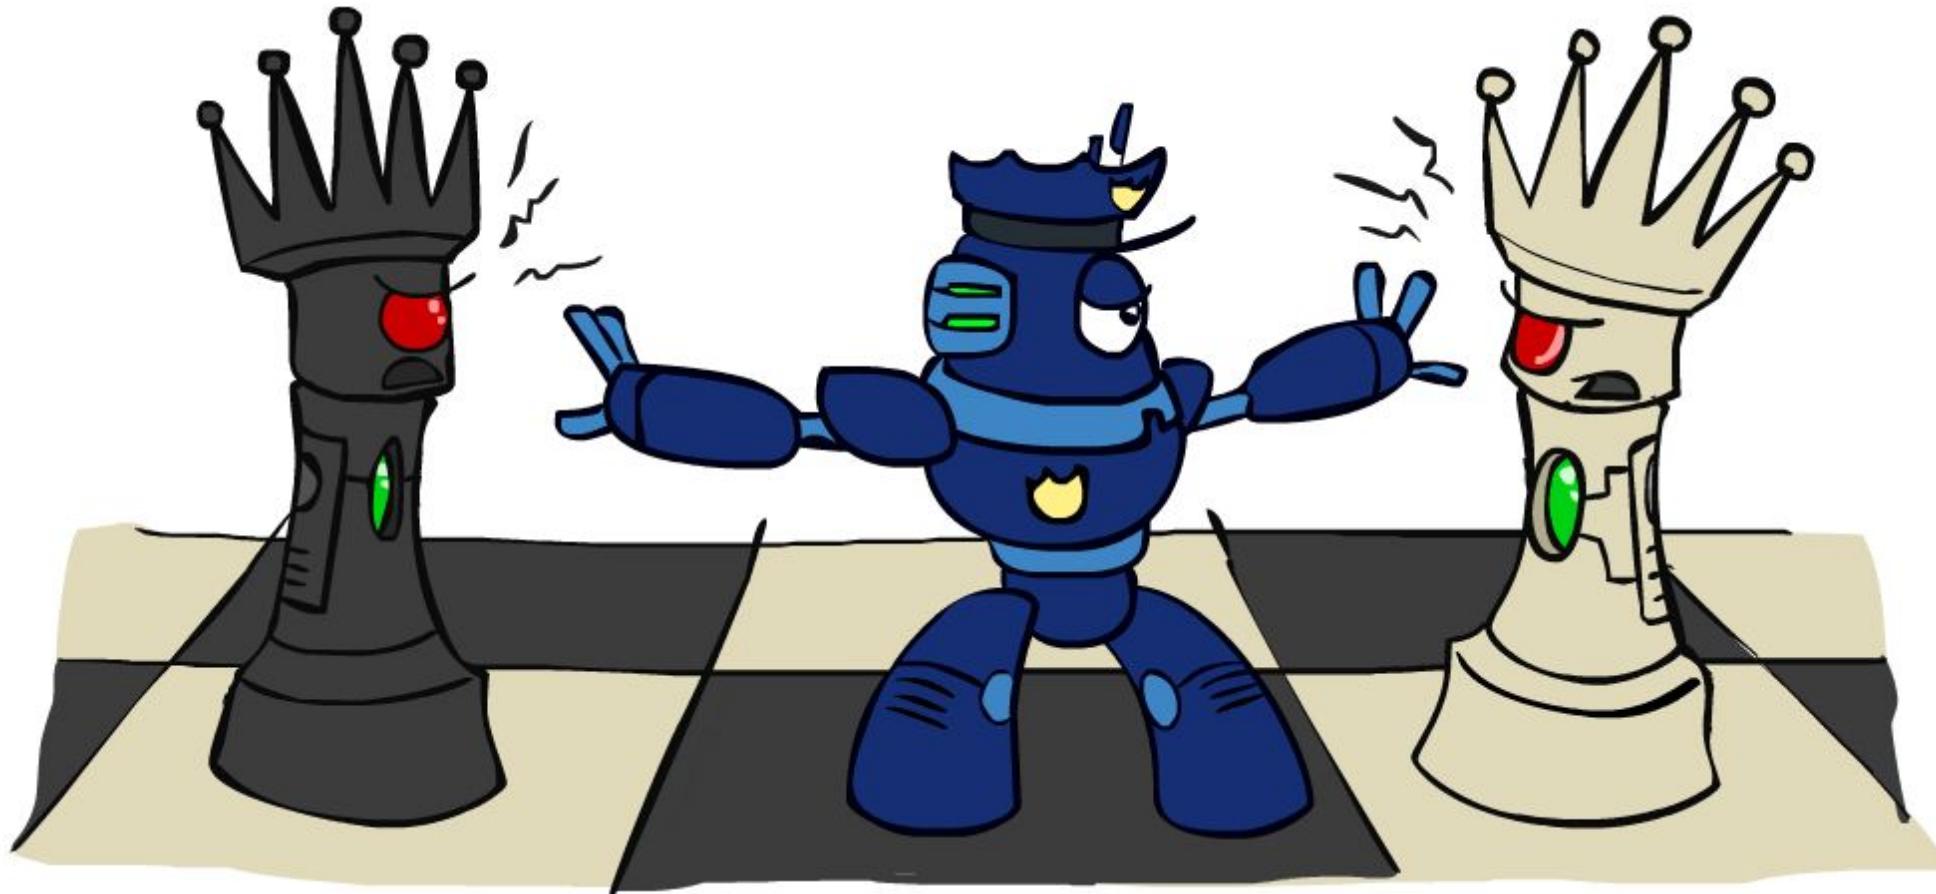


# Tree-Structured CSP



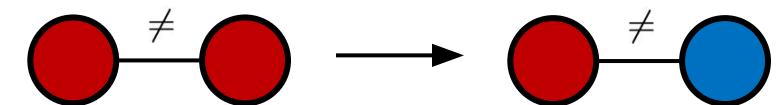
# Iterative Improvement

---

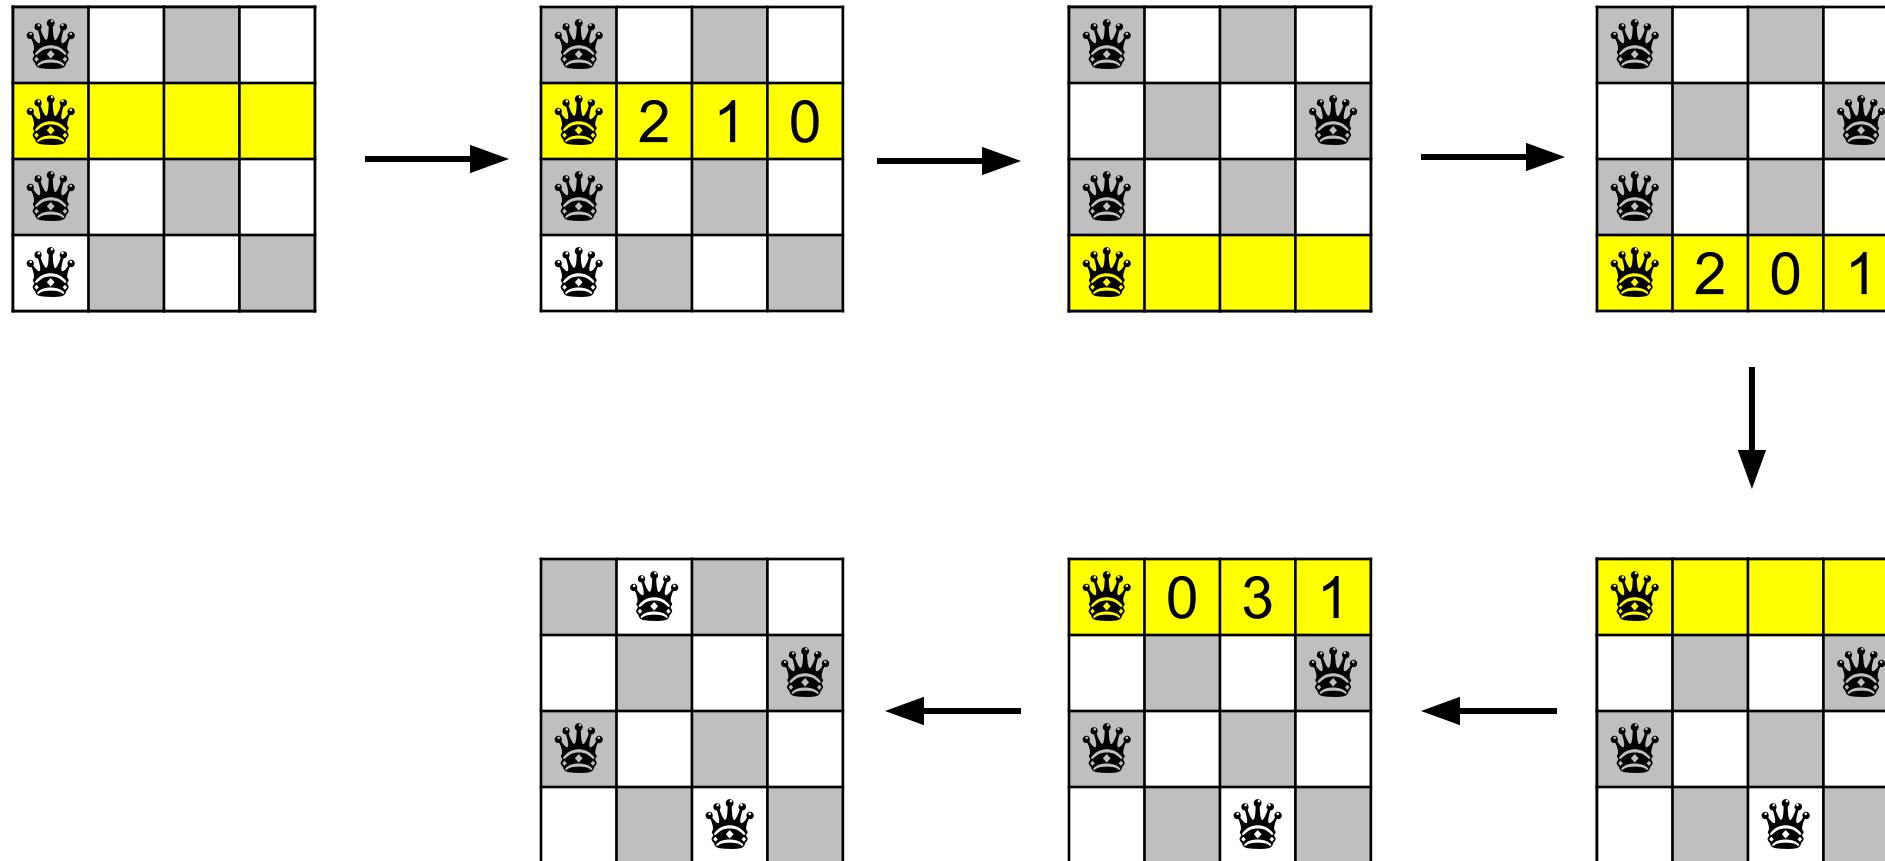


# Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
  - Take an assignment with unsatisfied constraints
  - Operators *reassign* variable values
  - No fringe! Live on the edge.
- Algorithm: While not solved,
  - Variable selection: randomly select any conflicted variable
  - Value selection: min-conflicts heuristic:
    - Choose a value that violates the fewest constraints
    - I.e., hill climb with  $h(n) = \text{total number of violated constraints}$



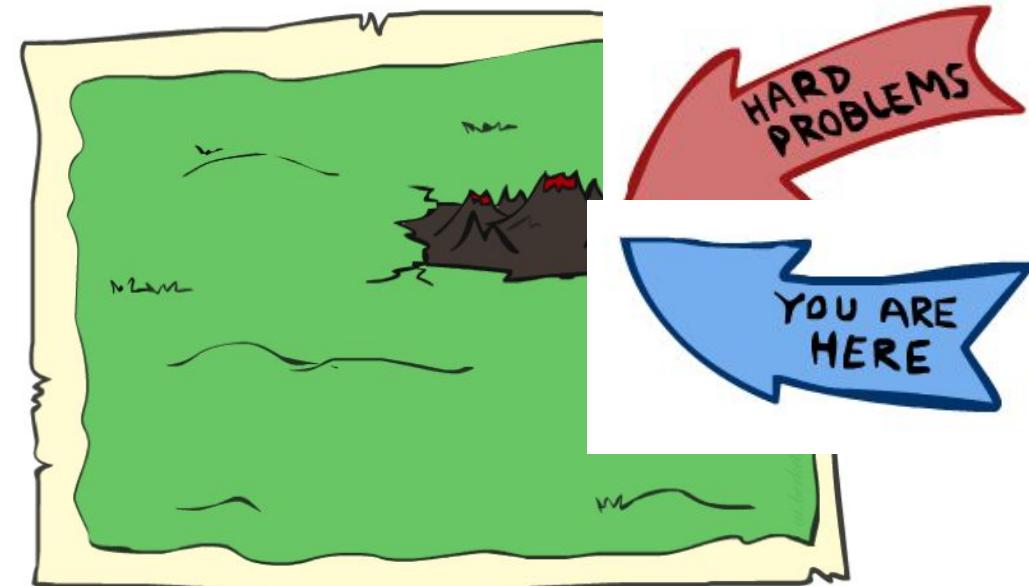
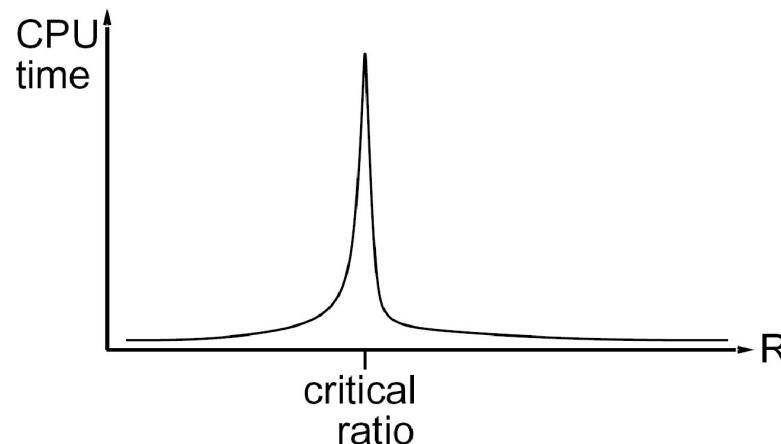
# Example: 4-Queens



# Performance of Min-Conflicts

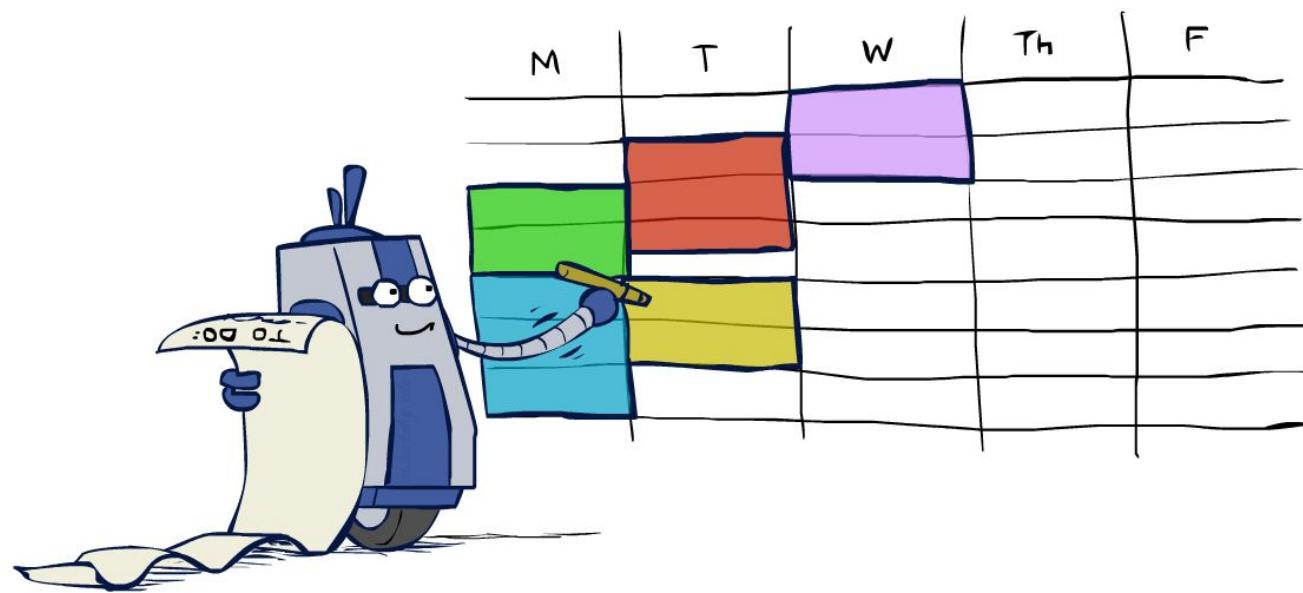
- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., n = 10,000,000)!
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$



# Summary: CSPs

- CSPs are a special kind of search problem:
  - States are partial assignments
  - Goal test defined by constraints
- Basic solution: backtracking search
- Speed-ups:
  - Ordering
  - Filtering
  - Structure
- Iterative min-conflicts is often effective in practice



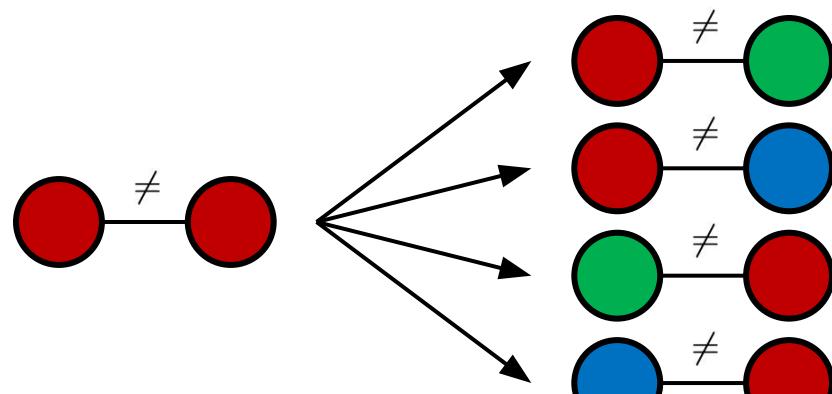
# Local Search

---



# Local Search

- Tree search keeps unexplored alternatives on the fringe (ensures completeness)
- Local search: improve a single option until you can't make it better
- New successor function: local changes



- Generally much faster and more memory efficient (but incomplete and suboptimal)
- Pretty much unavoidable when the state is “yourself”

# Hill Climbing

---

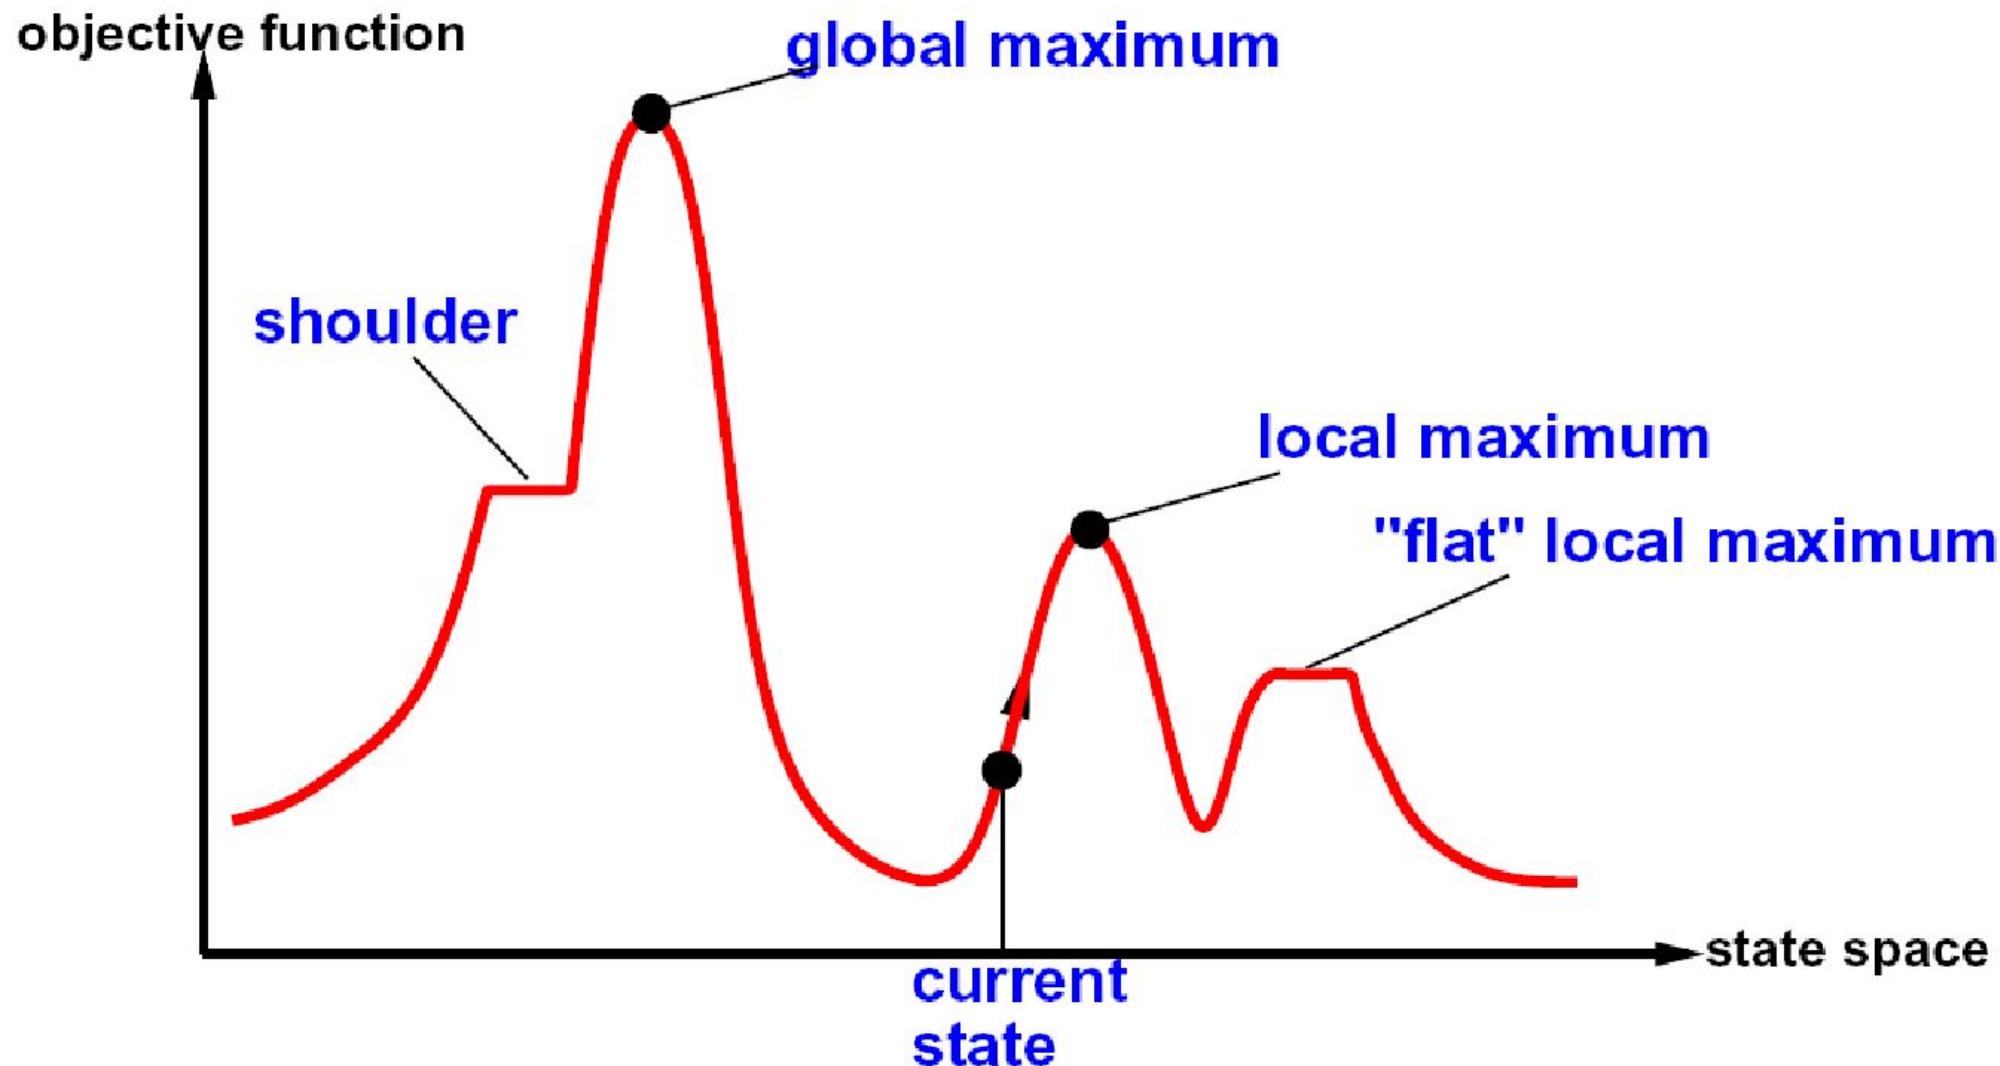
- Simple, general idea:
  - Start wherever
  - Repeat: move to the best neighboring state
  - If no neighbors better than current, quit



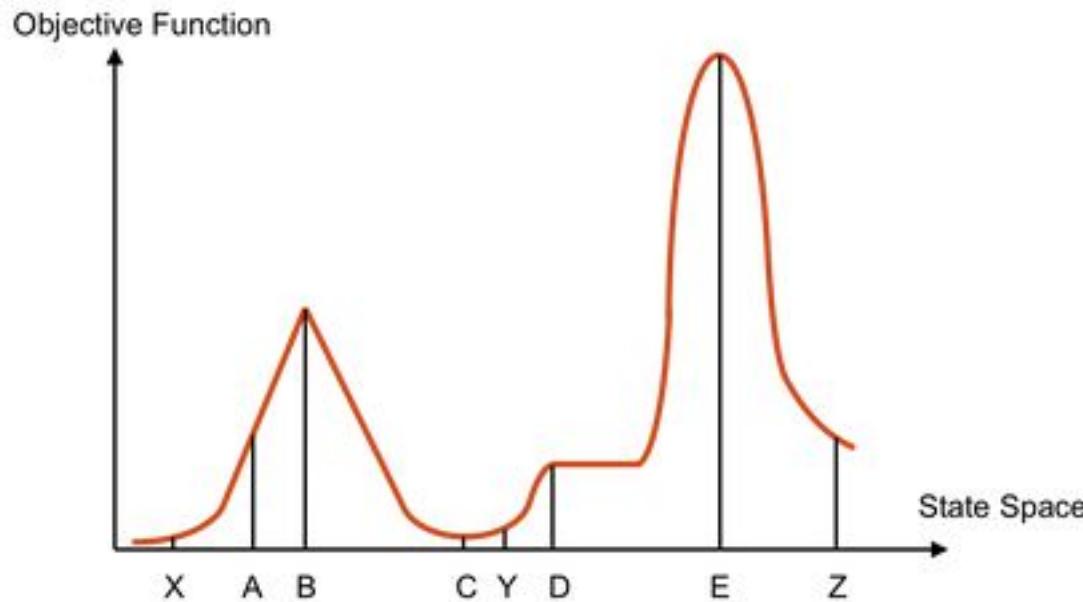
# Hill Climbing

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node
  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor  $\leftarrow$  a highest-valued successor of current
    if VALUE[neighbor]  $\leq$  VALUE[current] then return STATE[current]
    current  $\leftarrow$  neighbor
  end
```

# Hill Climbing Diagram



# Hill Climbing Quiz



Starting from X, where do you end up ?

Starting from Y, where do you end up ?

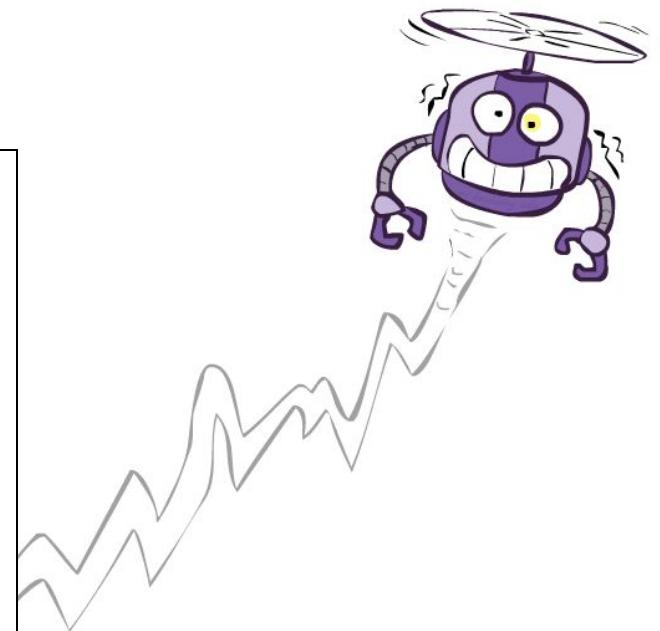
Starting from Z, where do you end up ?

# Simulated Annealing

- Idea: Escape local maxima by allowing downhill moves
  - But make them rarer as time goes on

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
          schedule, a mapping from time to “temperature”
  local variables: current, a node
                    next, a node
                    T, a “temperature” controlling prob. of downward steps

  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  for t  $\leftarrow$  1 to  $\infty$  do
    T  $\leftarrow$  schedule[t]
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  VALUE[next] – VALUE[current]
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
```



# Simulated Annealing

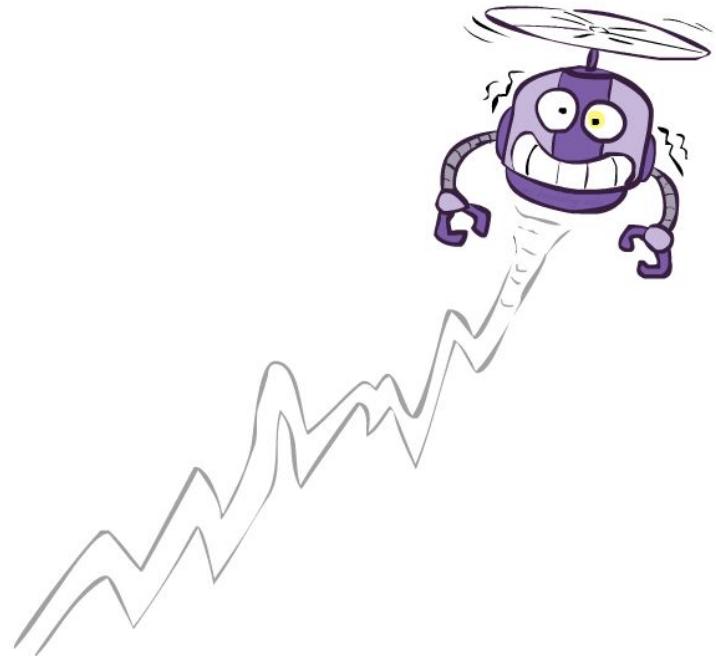
- Theoretical guarantee:

- Stationary distribution (Boltzmann):  $p(x) \propto e^{\frac{E(x)}{kT}}$
- If T decreased slowly enough,  
will converge to optimal state!

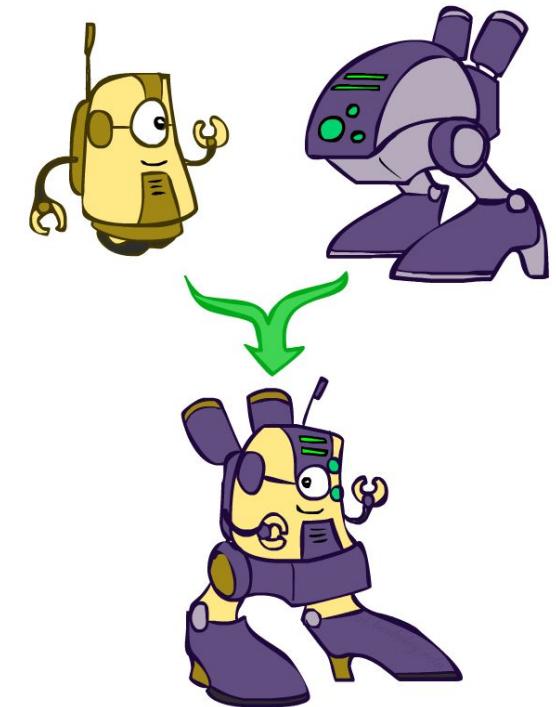
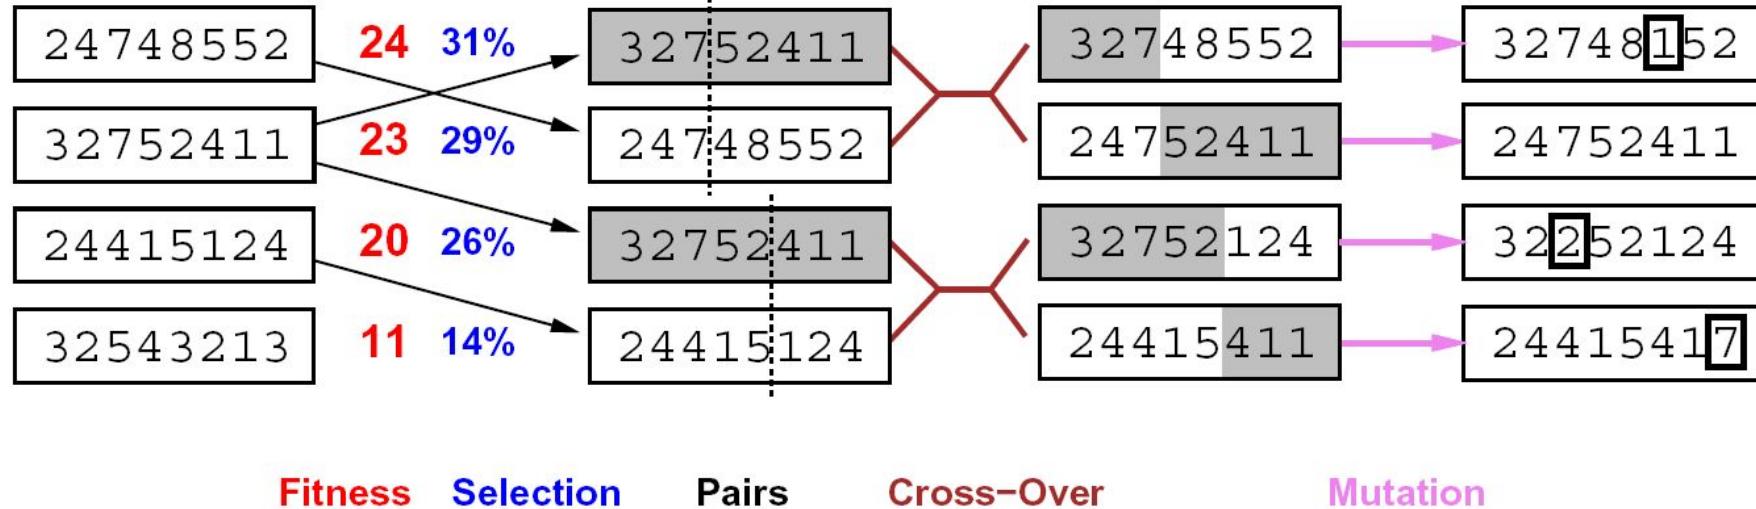
- Is this an interesting guarantee?

- Sounds like magic, but reality is reality:

- The more downhill steps you need to escape a local optimum,  
the less likely you are to ever make them all in a row
- “Slowly enough” may mean exponentially slowly
- Random restart hillclimbing also converges to optimal state...

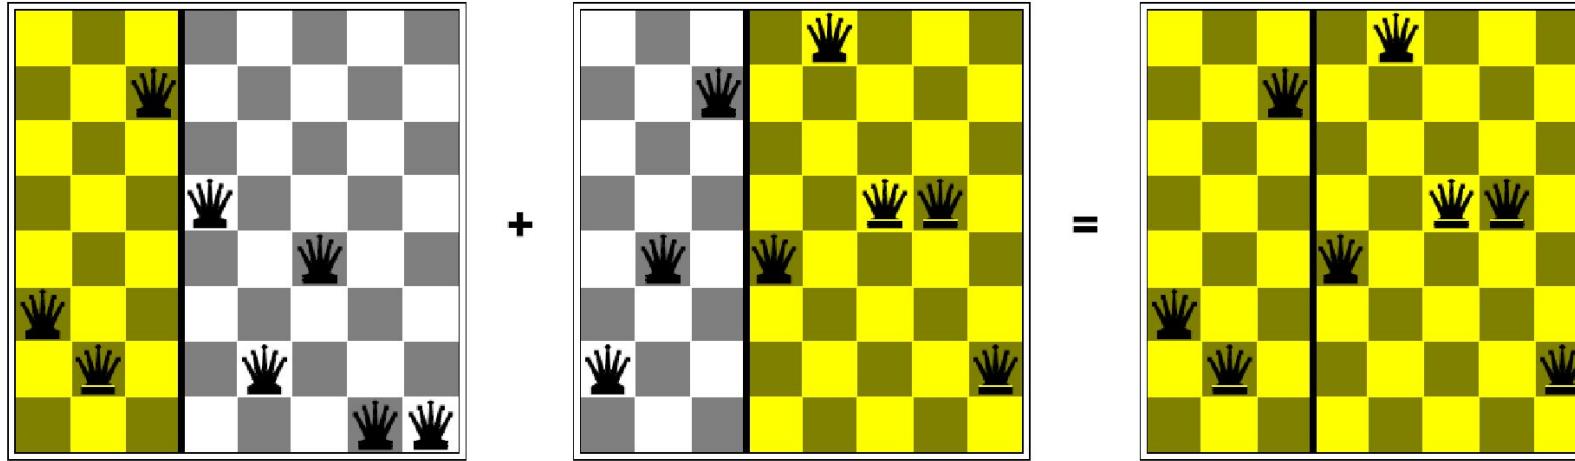


# Genetic Algorithms



- Genetic algorithms use a natural selection metaphor
  - Keep best N hypotheses at each step (selection) based on a fitness function
  - Also have pairwise crossover operators, with optional mutation to give variety
- Possibly the most misunderstood, misapplied (and even maligned) technique around

# Example: N-Queens



- Why does crossover make sense here?
- When wouldn't it make sense?
- What would mutation be?
- What would a good fitness function be?